

Stream processing with R and Amazon Kinesis

Budapest Data Forum 2016

Gergely Daroczi


@daroczig

June 16 2016





Contributions

Less  More

Contributions in the last year

3,477 total

Feb 9, 2015 – Feb 9, 2016

Longest streak

22 days

December 26 – January 16

Current streak

2 days

February 7 – February 8



Gergely Daróczy @daroczig · Apr 11


Just received my "I ♥ R" prepaid debit card from @CARD. Will be fun to use this #rstats designed card at #user2015 :)

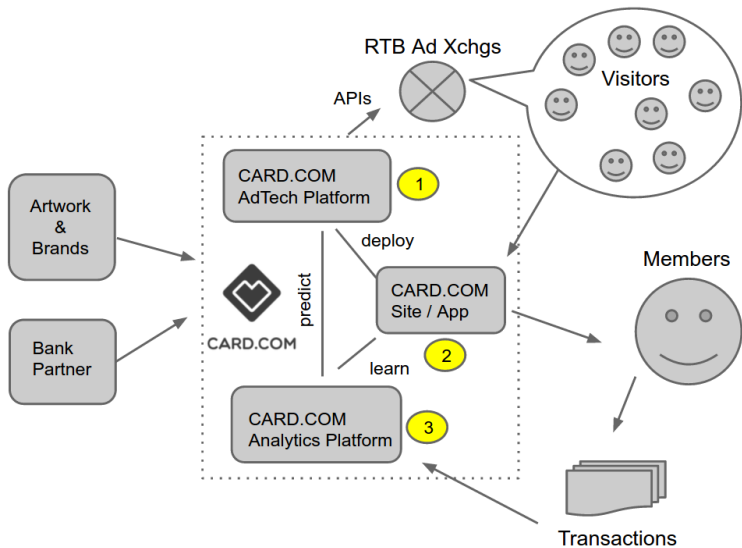


RETWEETS
10

FAVORITES
16







- card transaction processors
- card manufacturers
- CIP/KYC service providers
- online ad platforms
- remarketing networks
- licensing partners
- communication engines
- others



Classes and code fragments:

- Classes
- Code Snippets
- Input fields
 - Getting fields...please wait
- Info fields
 - Getting fields...please wait
- Output fields
 - Getting fields...please wait

Step name

Shorten Date

Class code

```

@ Processor R
import java.text.SimpleDateFormat;
import java.util.Date;
import java.text.ParseException;
import java.util.TimeZone;

private SimpleDateFormat df1 = new SimpleDateFormat("yyyy/MM/dd HH:mm:ss.SSS");
private SimpleDateFormat df2 = new SimpleDateFormat("yyyy-MM-dd HH*");

public boolean processRow(StepMetaInterface smi, StepDataInterface sdi) throws KettleException, ParseException
{
    Object[] r = getRow();
    if (r == null) {
        setOutputDone();
        return false;
    }

    if (first)
    {
        first = false;
    }

    // It is always safest to call createOutputRow() to ensure that your output row's Object[] is large
    // enough to handle any new fields you are creating in this step.
    r = createOutputRow(r, data.outputRowMeta.size());

    df2.setTimeZone(TimeZone.getTimeZone("America/Los_Angeles"));
        
```

Line #: 0

Fields Parameters Info steps Target steps

Fields

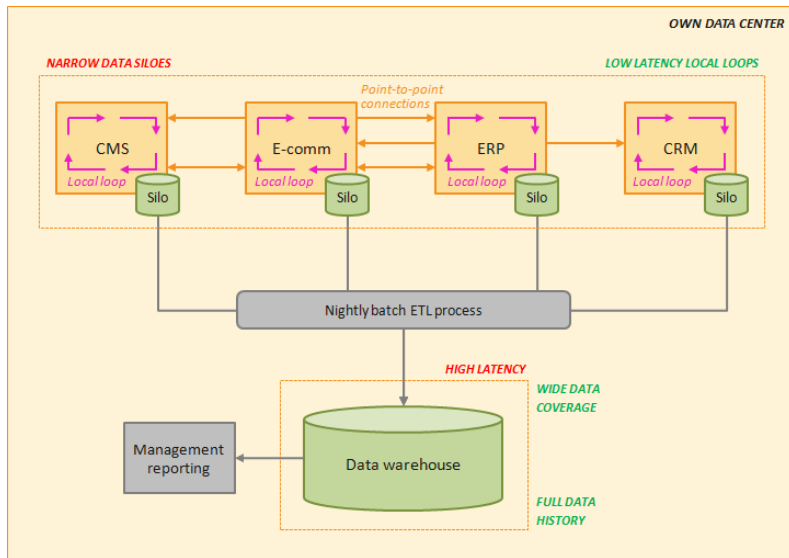
Clear the result fields?

#	Fieldname	Type	Length	Precision
1	RPT_DATE_SHORT	String		

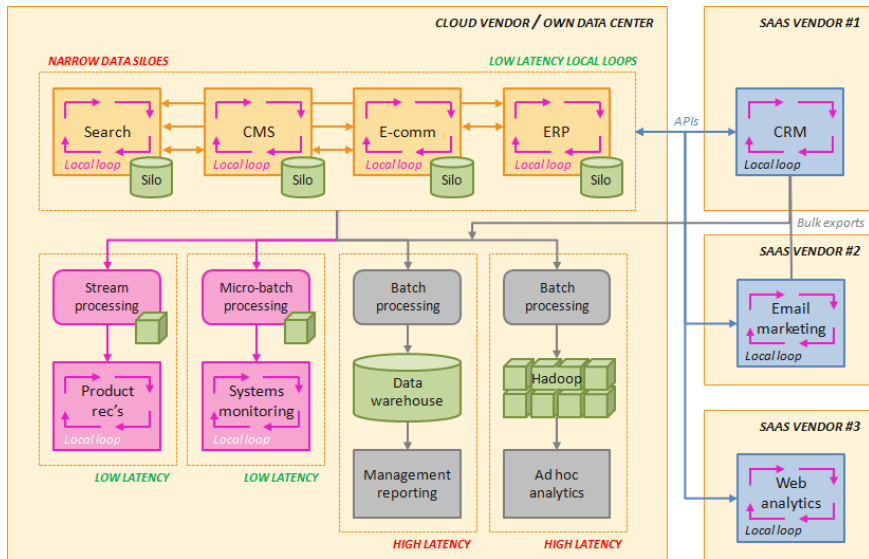
Help

OK Cancel Test class

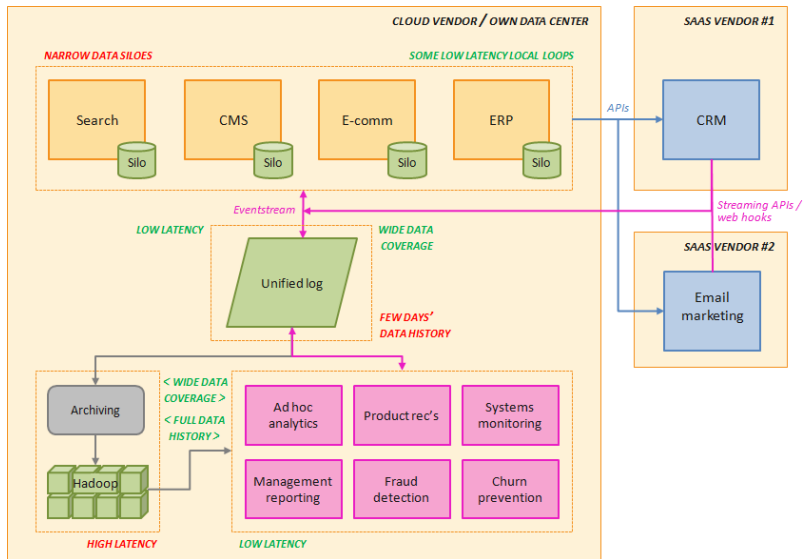





Source: SnowPlow



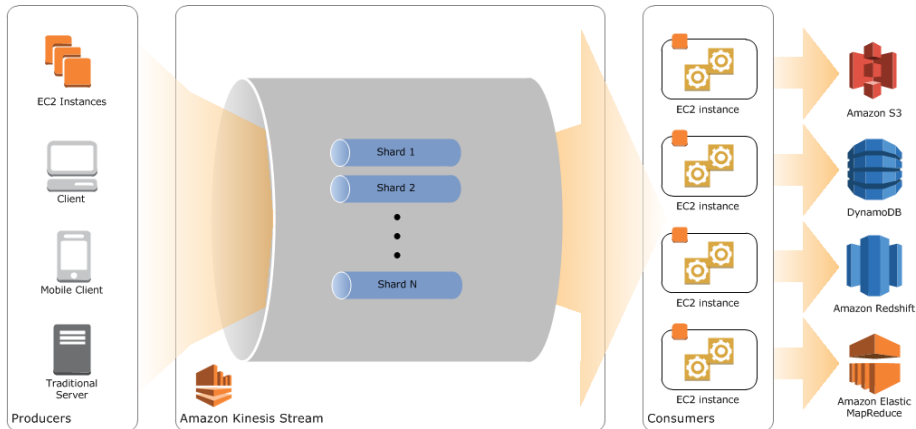
Source: SnowPlow



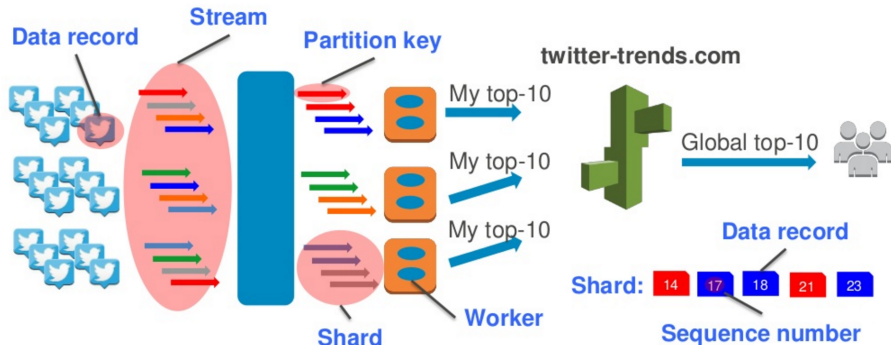
Source: SnowPlow



Source: [Kinesis Product Details](#)



Source: [Kinesis Developer Guide](#)



Source: AWS re:Invent 2013

Writing data to the stream:

- Amazon Kinesis Streams API (!)
- Amazon Kinesis Producer Library (KPL) from Java
- flume-kinesis
- Amazon Kinesis Agent

Reading data from the stream:

- Amazon Kinesis Streams API (!)
- Amazon Kinesis Client Library (KCL) from Java, Node.js, .NET, Python, Ruby

Managing streams:

- Amazon Kinesis Streams API (!)

Initialize connection to the SDK:

```
> library(rJava)
+ .jinit(classpath =
+       list.files(
+         '~/Projects/kineric/inst/java/',
+         full.names = TRUE),
+       parameters = "-Xmx512m")
```

What methods do we have there?

```
> kc <- .jnew('com.amazonaws.services.kinesis.AmazonKinesisClient')
> .jmethods(kc)
[1] "public void com.amazonaws.services.kinesis.AmazonKinesisClient.splitShard(java
[2] "public void com.amazonaws.services.kinesis.AmazonKinesisClient.splitShard(com
[3] "public void com.amazonaws.services.kinesis.AmazonKinesisClient.removeTagsFrom
[4] "public com.amazonaws.services.kinesis.model.ListStreamsResult com.amazonaws.s
[5] "public com.amazonaws.services.kinesis.model.ListStreamsResult com.amazonaws.s
[6] "public com.amazonaws.services.kinesis.model.ListStreamsResult com.amazonaws.s
[7] "public com.amazonaws.services.kinesis.model.ListStreamsResult com.amazonaws.s
[8] "public com.amazonaws.services.kinesis.model.GetShardIteratorResult com.amazon
[9] "public com.amazonaws.services.kinesis.model.GetShardIteratorResult com.amazon
[10] "public com.amazonaws.services.kinesis.model.GetShardIteratorResult com.amazon
```

Set API endpoint:

```
> kc$setEndpoint('kinesis.us-west-2.amazonaws.com', 'kinesis', 'us-west-2')
```

What streams do we have access to?

```
> kc$listStreams()
[1] "Java-Object[{StreamNames: [test_kinesis],HasMoreStreams: false}]"
> kc$listStreams()$getStreamNames()
[1] "Java-Object[[test_kinesis]]"
```

Describe this stream:

```
> (kcstream <- kc$describeStream(StreamName = 'test_kinesis')$getStreamDescription())
[1] "Java-Object[{StreamName: test_kinesis,StreamARN: arn:aws:kinesis:us-west-2:595
> kcstream$getStreamName()
[1] "test_kinesis"
> kcstream$getStreamStatus()
[1] "ACTIVE"
> kcstream$getShards()
[1] "Java-Object[{
{ShardId: shardId-0000000000000,HashKeyRange: {
StartingHashKey: 0,EndingHashKey: 17014118346046923173168730371588410572324},Sequen
{ShardId: shardId-00000000000001,HashKeyRange: {StartingHashKey: 17014118346046923173
```

```
> prr <- .jnew('com.amazonaws.services.kinesis.model.PutRecordRequest')
> prr$setStreamName('test_kinesis')
> prr$setData(J('java.nio.ByteBuffer')$wrap(.jbyte(charToRaw('foobar'))))
> prr$setPartitionKey('42')
> prr
[1] "Java-Object{{StreamName: test_kinesis,
      Data: java.nio.HeapByteBuffer[pos=0 lim=6 cap=6],
      PartitionKey: 42,}}"
> kc$putRecord(prr)
[1] "Java-Object{{ShardId: shardId-000000000001,
      SequenceNumber: 49562894160449444332153346371084313572324361665031176210}}"
```

First, we need a shared iterator, which expires in 5 minutes:

```
> sir <- .jnew('com.amazonaws.services.kinesis.model.GetShardIteratorRequest')
> sir$setStreamName('test_kinesis')
> sir$setShardId(.jnew('java/lang/String', '0'))
> sir$setShardIteratorType('TRIM_HORIZON')
> iterator <- kc$getShardIterator(sir)$getShardIterator()
```

Then we can use it to get records:

```
> gir <- .jnew('com.amazonaws.services.kinesis.model.GetRecordsRequest')
> gir$setShardIterator(iterator)
> kc$getRecords(gir)$getRecords()
[1] "Java-Object{[]}"
```

First, we need a shared iterator, which expires in 5 minutes:

```
> sir <- .jnew('com.amazonaws.services.kinesis.model.GetShardIteratorRequest')
> sir$setStreamName('test_kinesis')
> sir$setShardId(.jnew('java/lang/String', '0'))
> sir$setShardIteratorType('TRIM_HORIZON')
> iterator <- kc$getShardIterator(sir)$getShardIterator()
```

Then we can use it to get records:

```
> gir <- .jnew('com.amazonaws.services.kinesis.model.GetRecordsRequest')
> gir$setShardIterator(iterator)
> kc$getRecords(gir)$getRecords()
[1] "Java-Object{[]}"
```

The data was pushed to the other shard:

```
> sir$setShardId(.jnew('java/lang/String', '1'))
> iterator <- kc$getShardIterator(sir)$getShardIterator()
> gir$setShardIterator(iterator)
> kc$getRecords(gir)$getRecords()
[1] "Java-Object{[{SequenceNumber: 495628941604494443321533463710843135723243616650,
ApproximateArrivalTimestamp: Tue Jun 14 09:40:19 CEST 2016,
Data: java.nio.HeapByteBuffer[pos=0 lim=6 cap=6],
PartitionKey: 402777}]}"
```

```
> supply(kc$getRecords(gir)$getRecords(),
+       function(x)
+         rawToChar(x$getData()$array()))
[1] "foobar"
> iterator <- kc$getRecords(gir)$getNextShardIterator()
> gir$setShardIterator(iterator)
> kc$getRecords(gir)$getRecords()
[1] "Java-Object{ [] }"
```

Get the oldest available data again:

```
> iterator <- kc$getShardIterator(sir)$getShardIterator()
> gir$setShardIterator(iterator)
> kc$getRecords(gir)$getRecords()
[1] "Java-Object{ [{SequenceNumber: 49562894160449444332153346371084313572324361665031}] }"
```

Or we can refer to the last known sequence number as well:

```
> sir$setShardIteratorType('AT_SEQUENCE_NUMBER')
> sir$setStartingSequenceNumber('49562894160449444332153346371084313572324361665031')
> iterator <- kc$getShardIterator(sir)$getShardIterator()
> gir$setShardIterator(iterator)
> kc$getRecords(gir)$getRecords()
[1] "Java-Object{ [{SequenceNumber: 49562894160449444332153346371084313572324361665031}] }"
```

No need for both shards:

```
> ms <- .jnew('com.amazonaws.services.kinesis.model.MergeShardsRequest')
> ms$setShardToMerge('shardId-000000000000')
> ms$setAdjacentShardToMerge('shardId-000000000001')
> ms$setStreamName('test_kinesis')
> kc$mergeShards(ms)
```

What do we have now?

```
> kc$describeStream(StreamName = 'test_kinesis')$getStreamDescription()$getShards()
[1] "Java-Object{[
{ShardId: shardId-000000000000,HashKeyRange: {StartingHashKey: 0,EndingHashKey: 170
SequenceNumberRange: {
StartingSequenceNumber: 49562894160427143586954815717376297430913467927668719618,
EndingSequenceNumber: 49562894160438293959554081028945856364232263390243848194}},
{ShardId: shardId-000000000001,HashKeyRange: {StartingHashKey: 17014118346046923173
SequenceNumberRange: {
StartingSequenceNumber: 49562894160449444332153346340517833149186116289174700050,
EndingSequenceNumber: 49562894160460594704752611652087392082504911751749828626}},
{ShardId: shardId-000000000002,
ParentShardId: shardId-000000000000,
AdjacentParentShardId: shardId-000000000001,
HashKeyRange: {StartingHashKey: 0,EndingHashKey: 3402823669209384634633746074317682
SequenceNumberRange: {StartingSequenceNumber: 4956290499149767309970492434472701952
```

Ideas:

- R function managing/scaling shards and attaching R processors
- One R processor per shard (parallel threads on a node or cluster)
- Store started/finished sequence number in memory/DB (checkpointing)

Example (Design for a simple MVP on a single node)

mcparallel starts parallel R processes to evaluate the given expression

Ideas:

- R function managing/scaling shards and attaching R processors
- One R processor per shard (parallel threads on a node or cluster)
- Store started/finished sequence number in memory/DB (checkpointing)

Example (Design for a simple MVP on a single node)

mcparallel starts parallel R processes to evaluate the given expression

Problems:

- duplicated records after failure due to in-memory checkpoint
- needs a DB for running on a cluster
- hard to write good unit tests

Ideas:

- R function managing/scaling shards and attaching R processors
- One R processor per shard (parallel threads on a node or cluster)
- Store started/finished sequence number in memory/DB (checkpointing)

Example (Design for a simple MVP on a single node)

mcparallel starts parallel R processes to evaluate the given expression

Problems:

- duplicated records after failure due to in-memory checkpoint
- needs a DB for running on a cluster
- hard to write good unit tests

Why not use KCL then?

- An easy-to-use programming model for processing data

```
java -cp amazon-kinesis-client-1.6.1.jar \  
com.amazonaws.services.kinesis.multilang.MultiLangDaemon \  
test-kinesis.properties
```

- Scale-out and fault-tolerant processing (checkpointing via DynamoDB)
- Logging and metrics in CloudWatch
- The MultiLangDaemon spawns processes written in any language, communication happens via JSON messages sent over stdin/stdout
- Only a few events/methods to care about in the consumer application:
 - 1 initialize
 - 2 processRecords
 - 3 checkpoint
 - 4 shutdown

① initialize:

- Perform initialization steps
- Write “status” message to indicate you are done
- Begin reading line from STDIN to receive next action

② processRecords:

- Perform processing tasks (you may write a checkpoint message at any time)
- Write “status” message to STDOUT to indicate you are done.
- Begin reading line from STDIN to receive next action

③ shutdown:

- Perform shutdown tasks (you may write a checkpoint message at any time)
- Write “status” message to STDOUT to indicate you are done.
- Begin reading line from STDIN to receive next action

④ checkpoint:

- Decide whether to checkpoint again based on whether there is an error or not.

Step name

Shorten Date

Classes and code fragments:

Classes

Code Snippets

Input fields

Getting fields...please wait

Info fields

Getting fields...please wait

Output fields

Getting fields...please wait

Class code

```

@ Processor R
import java.text.SimpleDateFormat;
import java.util.Date;
import java.text.ParseException;
import java.util.TimeZone;

private SimpleDateFormat df1 = new SimpleDateFormat("yyyy/MM/dd HH:mm:ss.SSS");
private SimpleDateFormat df2 = new SimpleDateFormat("yyyy-MM-dd HH*");

public boolean processRow(StepMetaInterface smi, StepDataInterface sdi) throws KettleException, ParseException
{
    Object[] r = getRow();
    if (r == null) {
        setOutputDone();
        return false;
    }

    if (first)
    {
        first = false;
    }

    // It is always safest to call createOutputRow() to ensure that your output row's Object[] is large
    // enough to handle any new fields you are creating in this step.
    r = createOutputRow(r, data.outputRowMeta.size());

    df2.setTimeZone(TimeZone.getTimeZone("America/Los_Angeles"));

```

Line #: 0

Fields Parameters Info steps Target steps

Fields

Clear the result fields?

#	Fieldname	Type	Length	Precision
1	RPT_DATE_SHORT	String		

Help

OK Cancel Test class

```
#!/usr/bin/r -i

while (TRUE) {

  ## read and parse messages
  line <- fromJSON(readLines(n = 1))

  ## nothing to do unless we receive a record to process
  if (line$action == 'processRecords') {

    ## process each record
    lapply(line$records, function(r) {

      business_logic(r)
      cat(toJSON(list(action = 'checkpoint', checkpoint = r$sequenceNumber)))

    })
  }

  ## return response in JSON
  cat(toJSON(list(action = 'status', responseFor = line$action)))
}
```

