

Stream processing with R in AWS

AWR, AWR.KMS, AWR.Kinesis (R packages) used in ECS



CARD.COM

Gergely Daroczi

@daroczig



March 7, 2017







3,633 contributions in the year before last

Lead R Developer



3,470 contributions in the last year

Director of Analytics





Gergely Daróczki @daroczig · Apr 11

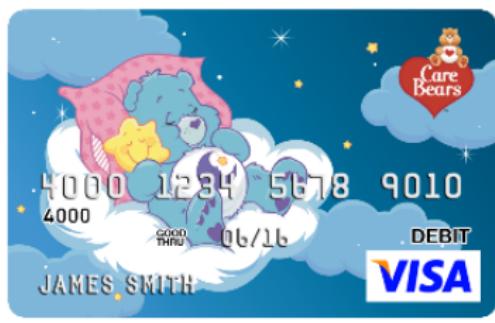
Just received my "I ❤️ R" prepaid debit card from @CARD. Will be fun to use this #rstats designed card at #user2015 :)

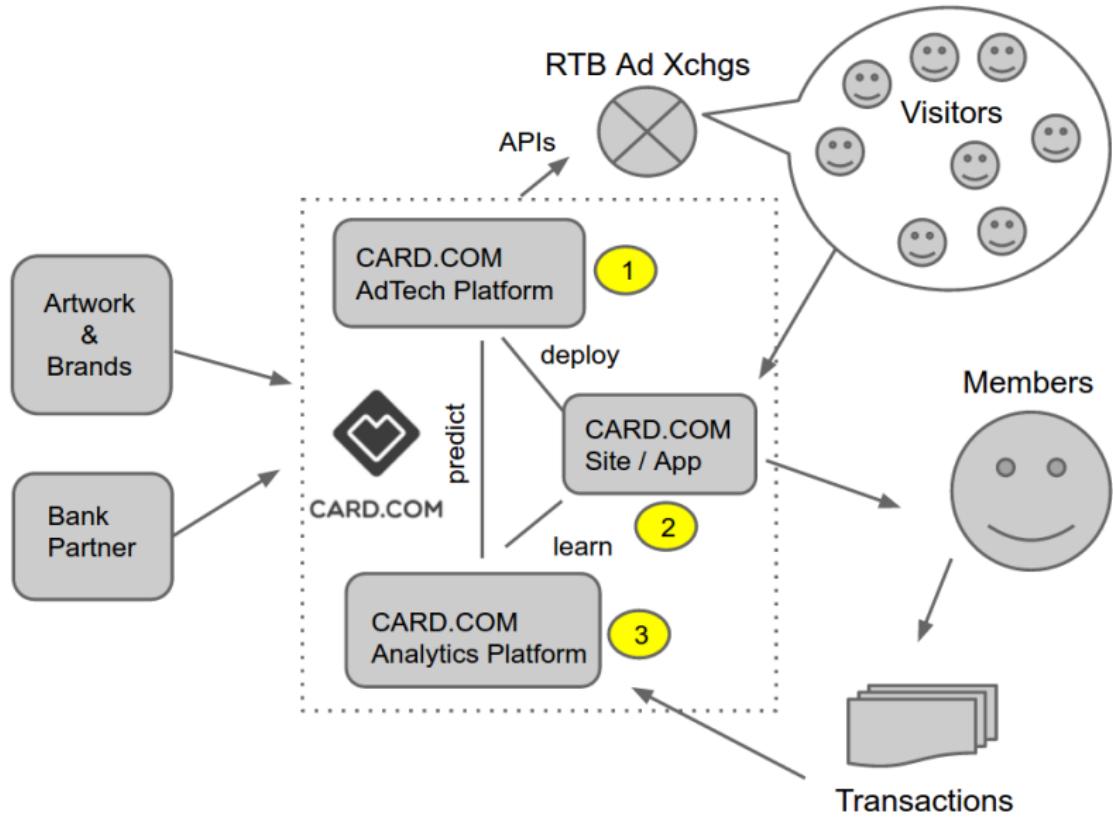


RETWEETS
10

FAVORITES
16







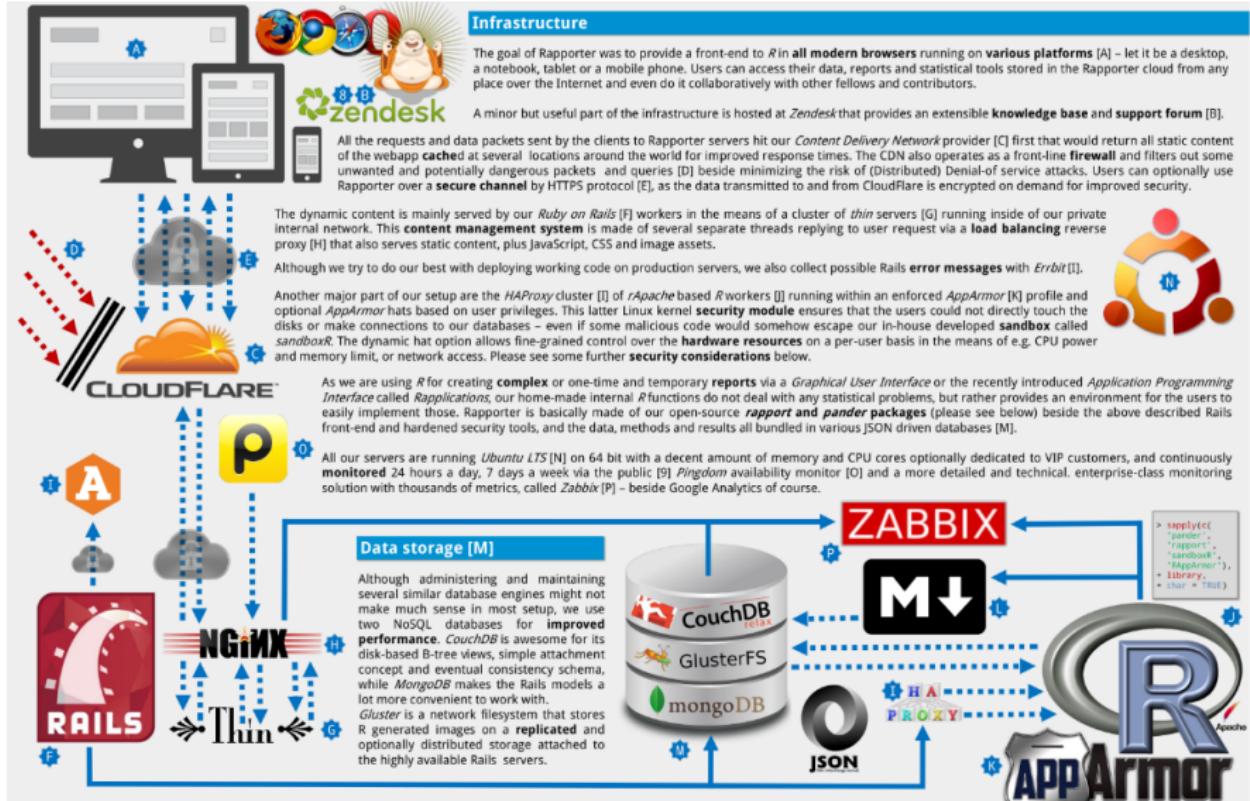
- card transaction processors
- card manufacturers
- CIP/KYC service providers
- online ad platforms
- remarketing networks
- licensing partners
- communication engines
- others



Why not Hadoop instead of MySQL?



CARD.COM





Why R?



User Defined Java Class

Step name: ShortenDate

Classes and code fragments:

- Classes
- Code Snippets
- Input fields
- Getting fields...please wait
- Info fields
- Getting fields...please wait
- Output fields
- Getting fields...please wait

Class code

```
Processor
import java.text.SimpleDateFormat;
import java.util.Date;
import java.text.ParseException;
import java.util.TimeZone;

private SimpleDateFormat df1 = new SimpleDateFormat("yyyy/MM/dd HH:mm:ss.SSS");

private SimpleDateFormat df2 = new SimpleDateFormat("yyyy-MM-dd HH");

public boolean processRow(StepMetaInterface smi, StepDataInterface sdi) throws KettleException, ParseException
{
    Object[] r = getRow();
    if (r == null) {
        setOutputDone();
        return false;
    }

    if (first)
    {
        first = false;
    }

    // It is always safest to call createOutputRow() to ensure that your output row's Object[] is large
    // enough to handle any new fields you are creating in this step.
    r = createOutputRow(r, data.outputRowMeta.size());

    df2.setTimeZone(TimeZone.getTimeZone("America/Los_Angeles"));
}
```

Line #: 0

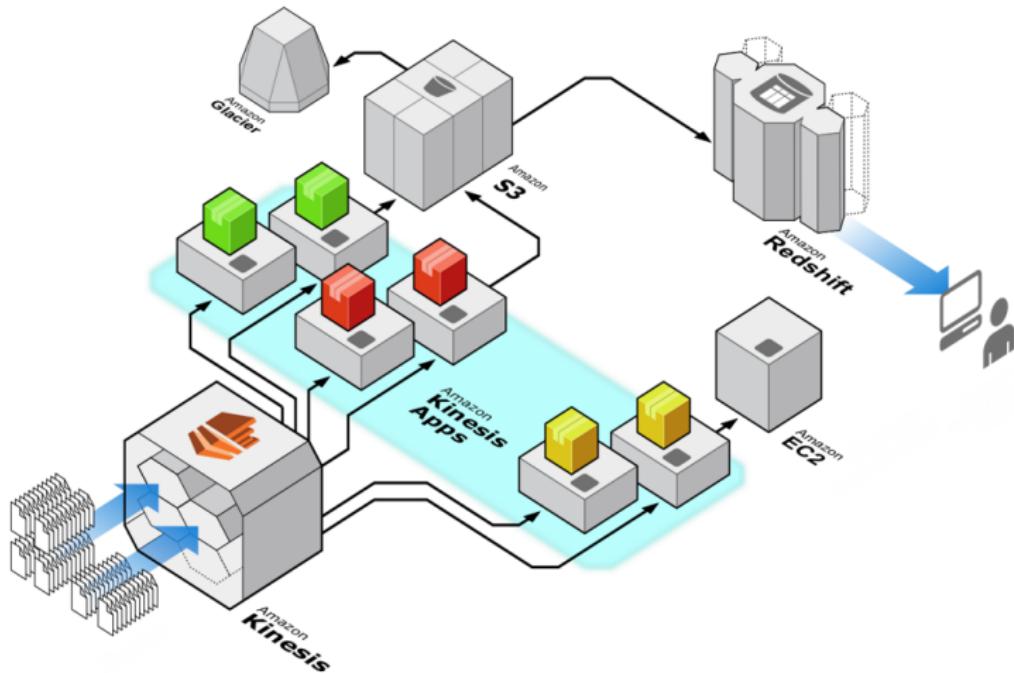
Fields Parameters Info steps Target steps

Fields Clear the result fields?

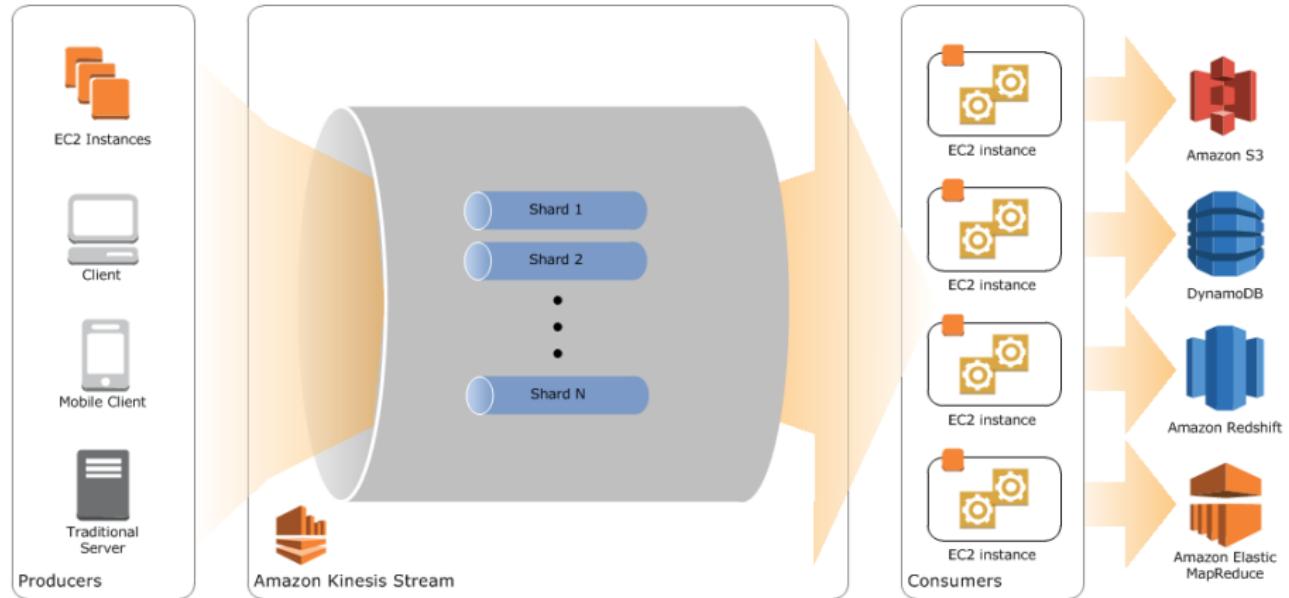
#	Fieldname	Type	Length	Precision
1	RPT_DATE_SHORT	String		

Help OK Cancel Test class

Why Amazon Kinesis?

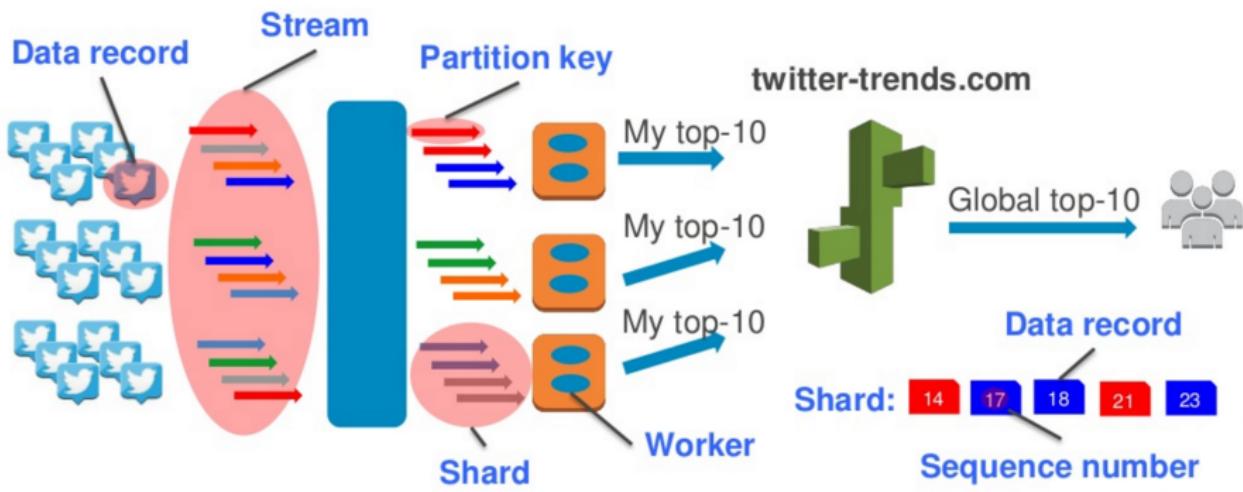


Source: Kinesis Product Details

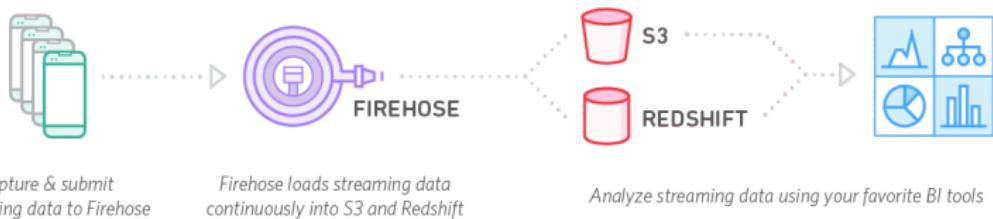


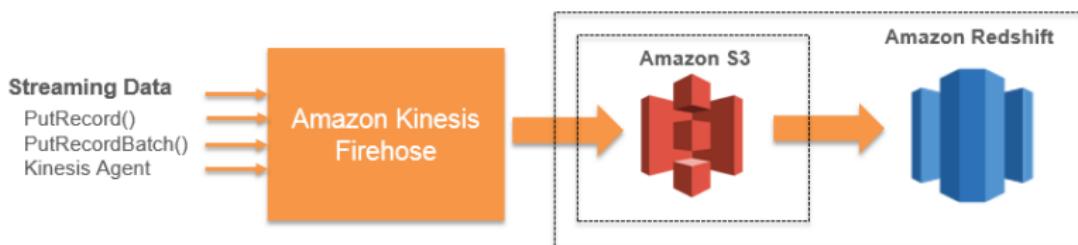
Source: Kinesis Developer Guide

Intro to Amazon Kinesis Shards



Source: AWS re:Invent 2013





The S3 Object System



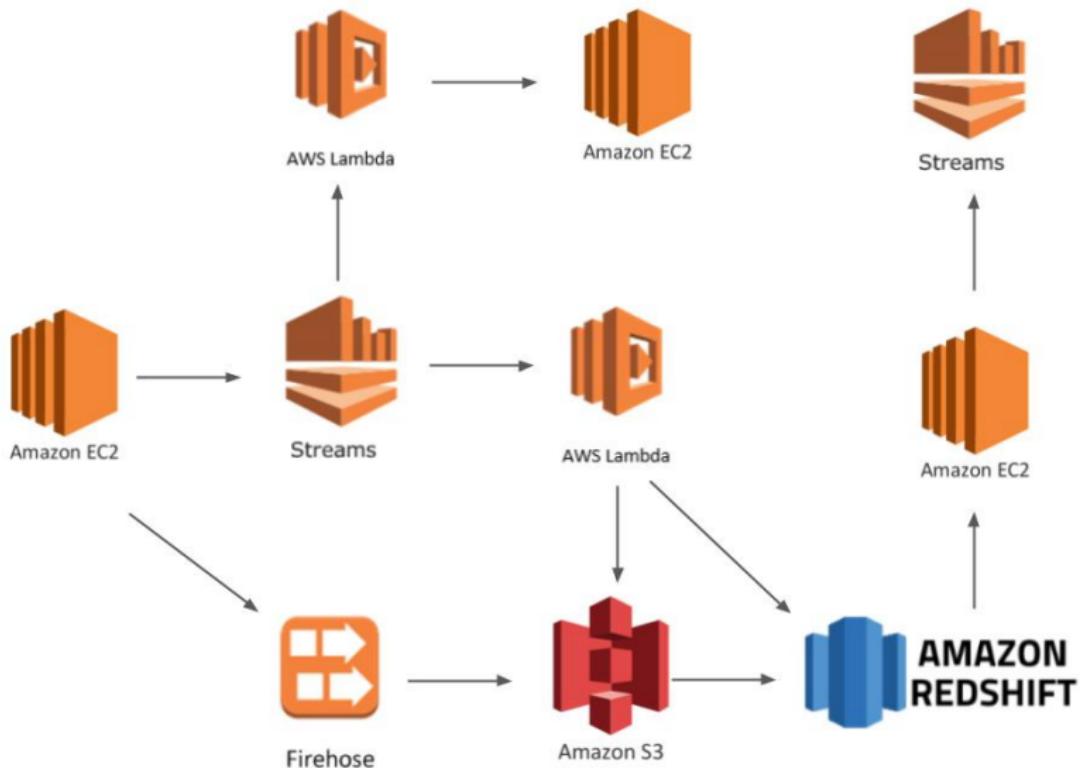
CARD.COM

```
> x <- 3.14
> attr(x, 'class') <- 'standard'

> print.standard <- function(x, ...) {
+     ## SLA
+     if (runif(1) * 100 > 99.9) {
+         Sys.sleep(20)
+     }
+     futile.logger::flog.info(x)
+ }
```

```
> while (TRUE) print(x)
INFO [2017-03-03 22:27:57] 3.14
INFO [2017-03-03 22:27:57] 3.14
INFO [2017-03-03 22:27:57] 3.14
INFO [2017-03-03 22:28:17] 3.14
INFO [2017-03-03 22:28:17] 3.14
```

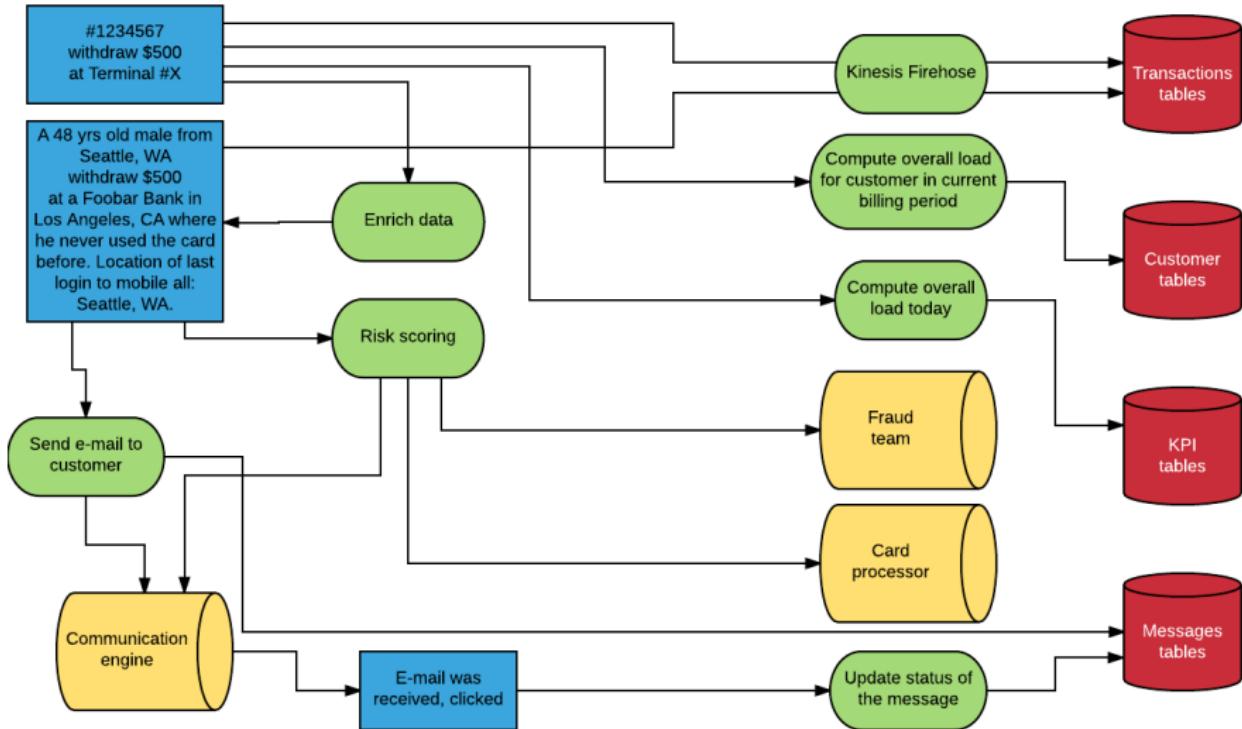
S4: Multiple Dispatch



Example use-case



CARD.COM



Writing data to the stream:

- Amazon Kinesis Streams API, SDK
- Amazon Kinesis Producer Library (KPL) from Java
- flume-kinesis
- Amazon Kinesis Agent

Reading data from the stream:

- Amazon Kinesis Streams API, SDK
- Amazon Kinesis Client Library (KCL) from Java, Node.js, .NET, Python, Ruby

Managing streams:

- Amazon Kinesis Streams API (!)

Now We Need an R Client!



CARD.COM

```
> library(rJava)
> .jinit(classpath = list.files('~/Projects/AWR/inst/java/', full.names = TRUE))

> kc <- .jnew('com.amazonaws.services.kinesis.AmazonKinesisClient')
> kc$setEndpoint('kinesis.us-west-2.amazonaws.com', 'kinesis', 'us-west-2')

> sir <- .jnew('com.amazonaws.services.kinesis.model.GetShardIteratorRequest')
> sir$setStreamName('test_kinesis')
> sir$setShardId(.jnew('java/lang/String', '0'))
> sir$setShardIteratorType('TRIM_HORIZON')
> iterator <- kc$getShardIterator(sir)$getShardIterator()

> grr <- .jnew('com.amazonaws.services.kinesis.model.GetRecordsRequest')
> grr$setShardIterator(iterator)
> kc$getRecords(grr)$getRecords()
[1] "Java-Object{[SequenceNumber: 495628941604494443321533463710843135723243616650,
ApproximateArrivalTimestamp: Tue Jun 14 09:40:19 CEST 2016,
Data: java.nio.HeapByteBuffer[pos=0 lim=6 cap=6],PartitionKey: 42]}"

> sapply(kc$getRecords(grr)$getRecords(),
+         function(x)
+             rawToChar(x$getData()$array()))
[1] "foobar"
```

Let's merge two shards:

```
> ms <- .jnew('com.amazonaws.services.kinesis.model.MergeShardsRequest')
> ms$setShardToMerge('shardId-000000000000')
> ms$setAdjacentShardToMerge('shardId-000000000001')
> ms$setStreamName('test_kinesis')
> kc$mergeShards(ms)
```

What do we have now?

```
> kc$describeStream(StreamName = 'test_kinesis')$getStreamDescription()$getShards()
[1] "Java-Object{[
{ShardId: shardId-000000000000,HashKeyRange: {StartingHashKey: 0,EndingHashKey: 170
SequenceNumberRange: {
StartingSequenceNumber: 49562894160427143586954815717376297430913467927668719618,
EndingSequenceNumber: 49562894160438293959554081028945856364232263390243848194}},
{ShardId: shardId-000000000001,HashKeyRange: {StartingHashKey: 17014118346046923173
SequenceNumberRange: {
StartingSequenceNumber: 4956289416044944332153346340517833149186116289174700050,
EndingSequenceNumber: 49562894160460594704752611652087392082504911751749828626}},
{ShardId: shardId-000000000002,
ParentShardId: shardId-000000000000,
AdjacentParentShardId: shardId-000000000001,
HashKeyRange: {StartingHashKey: 0,EndingHashKey: 3402823669209384634633746074317682
SequenceNumberRange: {StartingSequenceNumber: 495629049914976730997049243472701952
Gergely Daroczi (@daroczig) Stream processing using AWR github.com/cardcorp/AWR 22 / 71
```



- An *easy-to-use* programming model for processing data

```
java -cp amazon-kinesis-client-1.7.3.jar \
com.amazonaws.services.kinesis.multilang.MultiLangDaemon \
app.properties
```

- *Scalable* and *fault-tolerant* processing (checkpointing via DynamoDB)
- Logging and metrics in CloudWatch
- The **MultiLangDaemon** spawns processes written in any language, communication happens via JSON messages sent over stdin/stdout
- Only a few events/methods to care about in the consumer application:
 - ① initialize
 - ② processRecords
 - ③ checkpoint
 - ④ shutdown



① initialize:

- Perform initialization steps
- Write “status” message to indicate you are done
- Begin reading line from STDIN to receive next action

② processRecords:

- Perform processing tasks (you may write a checkpoint message at any time)
- Write “status” message to STDOUT to indicate you are done.
- Begin reading line from STDIN to receive next action

③ shutdown:

- Perform shutdown tasks (you may write a checkpoint message at any time)
- Write “status” message to STDOUT to indicate you are done.
- Begin reading line from STDIN to receive next action

④ checkpoint:

- Decide whether to checkpoint again based on whether there is an error or not.

Again: Why R?

User Defined Java Class

Step name: ShortenDate

Classes and code fragments:

- Clases
- Code Snippets
- Input fields
- Getting fields...please wait**
- Info fields
- Getting fields...please wait**
- Output fields
- Getting fields...please wait**

Class code

```

Processor
import java.text.SimpleDateFormat;
import java.util.Date;
import java.text.ParseException;
import java.util.TimeZone;

private SimpleDateFormat df1 = new SimpleDateFormat("yyyy/MM/dd HH:mm:ss.SSS");

private SimpleDateFormat df2 = new SimpleDateFormat("yyyy-MM-dd HH");

public boolean processRow(StepMetaInterface smi, StepDataInterface sdi) throws KettleException, ParseException
{
    Object[] r = getRow();
    if (r == null) {
        setOutputDone();
        return false;
    }

    if (first)
    {
        first = false;
    }

    // It is always safest to call createOutputRow() to ensure that your output row's Object[] is large
    // enough to handle any new fields you are creating in this step.
    r = createOutputRow(r, data.outputRowMeta.size());

    df2.setTimeZone(TimeZone.getTimeZone("America/Los_Angeles"));
}

```

Line #: 0

Fields Parameters Info steps Target steps

#	Fieldname	Type	Length	Precision
1	RPT_DATE_SHORT	String		

Clear the result fields?

 Help OK Cancel Test class



```
#!/usr/bin/r -i

while (TRUE) {

    ## read and parse JSON messages
    line <- fromJSON(readLines(n = 1))

    ## nothing to do unless we receive records to process
    if (line$action == 'processRecords') {

        ## process each record
        lapply(line$records, function(r) {

            business_logic(fromJSON(rawToChar(base64_dec(r$data))))
            cat(toJSON(list(action = 'checkpoint', checkpoint = r$sequenceNumber)))

        })
    }

    ## return response in JSON
    cat(toJSON(list(action = 'status', responseFor = line$action)))
}

}
```



```
#!/usr/bin/r -i

while (TRUE) {

    ## read and parse JSON messages
    line <- fromJSON(readLines(n = 1))

    ## nothing to do unless we receive records to process
    if (line$action == 'processRecords') {

        ## process each record
        lapply(line$records, function(r) {

            business_logic(fromJSON(rawToChar(base64_dec(r$data))))
            cat(toJSON(list(action = 'checkpoint', checkpoint = r$sequenceNumber)))

        })
    }

    ## return response in JSON
    cat(toJSON(list(action = 'status', responseFor = line$action)))
}

}
```

```
> install.packages('AWR.Kinesis')
also installing the dependency 'AWR'

trying URL 'https://cloud.r-project.org/src/contrib/AWR_1.11.89.tar.gz'
Content type 'application/x-gzip' length 3125 bytes

trying URL 'https://cloud.r-project.org/src/contrib/AWR.Kinesis_1.7.3.tar.gz'
Content type 'application/x-gzip' length 3091459 bytes (2.9 MB)

* installing *source* package 'AWR' ...
** testing if installed package can be loaded
trying URL 'https://gitlab.com/cardcorp/AWR/repository/archive.zip?ref=1.11.89'
downloaded 58.9 MB
* DONE (AWR)

* installing *source* package 'AWR.Kinesis' ...
* DONE (AWR.Kinesis)
```

Business logic coded in R (demo_app.R):

```
library(AWR.Kinesis)
kinesis_consumer(processRecords = function(records) {
  flog.info(jsonlite::toJSON(records))
})
```

Business logic coded in R (demo_app.R):

```
library(AWR.Kinesis)
kinesis_consumer(processRecords = function(records) {
  flog.info(jsonlite::toJSON(records))
})
```

Note

This is not something you should run in RStudio.

Business logic coded in R (demo_app.R):

```
library(AWR.Kinesis)
kinesis_consumer(processRecords = function(records) {
  flog.info(jsonlite::toJSON(records))
})
```

Config file for the MultiLangDaemon (demo_app.properties):

```
executableName = ./demo_app.R
streamName = demo_stream
applicationName = demo_app
```

Start the MultiLangDaemon:

```
/usr/bin/java -cp AWR/java/*:AWR.Kinesis/java/*:./ \
  com.amazonaws.services.kinesis.multilang.MultiLangDaemon \
  ./demo_app.properties
```

```
library(futile.logger)
library(AWR.Kinesis)

kinesis_consumer(
  initialize      = function()
    flog.info('Hello'),
  processRecords = function(records)
    flog.info(paste('Received', nrow(records), 'records from Kinesis')),
  shutdown        = function()
    flog.info('Bye'),
  updater         = list(
    list(1, function()
      flog.info('Updating some data every minute')),
    list(1/60*10, function()
      flog.info(paste(
        'This is a high frequency updater call',
        'running every 10 seconds')))),
  checkpointing = 1,
  logfile = '/logs/logger.log')
```

Note

In theory you could, but this is not something you should run in RStudio.

- ① Create a Kinesis Stream
- ② Create an IAM user with DynamoDB and Kinesis permissions
- ③ Write data to the Stream
- ④ Run the MultiLangDaemon referencing the properties file

Note

In theory you could, but this is not something you should run in RStudio.

- ① Create a Kinesis Stream
- ② Create an IAM user with DynamoDB and Kinesis permissions
- ③ Write data to the Stream
- ④ Run the MultiLangDaemon referencing the properties file





S Gergely Daróczig

Amazon Kinesis Stream

https://console.aws.amazon.com/kinesis/home?region=us-east-1#/

Services Resource Groups Gergely Daróczig N. Virginia Support

Amazon Kinesis

Amazon Kinesis services make it easier to work with real-time streaming data in the AWS Cloud.

 Amazon Kinesis Firehose

Continuously deliver streaming data to Amazon S3, Amazon Redshift, and Amazon Elasticsearch Service.

[Go to Firehose](#)

[Learn more about Firehose](#)

 Amazon Kinesis Analytics

Analyze streaming data from Amazon Kinesis Firehose and Amazon Kinesis Streams in real-time using SQL.

[Go to Analytics](#)

[Learn more about Analytics](#)

 Amazon Kinesis Streams

Collect and stream data for ordered, replayable, real-time processing.

[Go to Streams](#)

[Learn more about Streams](#)

Amazon Kinesis documentation and support

[Firehose documentation](#) | [Analytics documentation](#) | [Streams documentation](#) | [Forums](#)

Feedback English © 2008 - 2017, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

Create a Kinesis Stream



Amazon Kinesis Stream x https://console.aws.amazon.com/kinesis/home?region=us-east-1#/streams/create

Services Resource Groups Gergely Darócz N. Virginia Support

Create stream

Stream name* test-AWR

Streams Firehose Analytics

Shards

A shard is a unit of throughput capacity. Each shard ingests up to 1MB/sec and 1000 records/sec, and emits up to 2MB/sec. To accommodate for higher or lower throughput, the number of shards can be modified after the stream is created using the API. [Learn more](#)

Producers → Stream → Consumers

Estimate the number of shards you'll need

Number of shards* 1

You can provision up to 48 more shards before hitting your account limit of 50. [Learn more or request a shard limit increase for this account](#)

Total stream capacity Values are calculated based on the number of shards entered above.

Write 1 MB per second

1000 Records per second

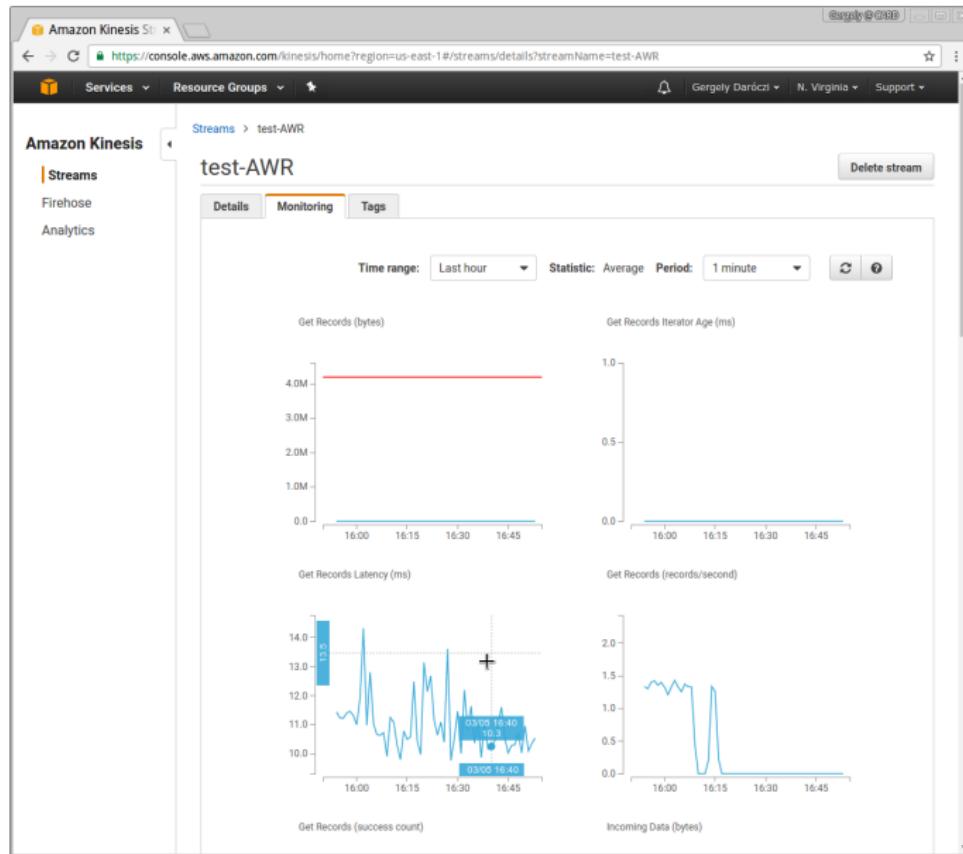
Read 2 MB per second

* Required Cancel Create stream

Check the Kinesis Stream



CARD.COM





Create an IAM user

IAM Management

https://console.aws.amazon.com/iam/home?region=us-east-1#users/AWR

Services Resource Groups

Gergely Daróczi Global Support

Search IAM

Users > AWR

Users: AWR

User ARN: arn:aws:iam::█████████████████████:user/AWR

Path: /

Creation time: 2017-02-13 16:04 PST

Permissions Groups (0) Security credentials Access Advisor

Add permissions Number of attached policies 5

- AmazonEC2ContainerRegistryFullAccess - AWS Managed policy
- CloudWatchFullAccess - AWS Managed policy
- cloudwatch-putmetrics - Managed policy
- AmazonDynamoDBFullAccess - AWS Managed policy
- AmazonKinesisFullAccess - AWS Managed policy

Add inline policy

Feedback English

© 2008-2017, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use



```
library(rJava)
.jcall("java/lang/System", "S", "setProperty", "aws.profile", "personal")

library(AWR.Kinesis)
library(jsonlite)
library(futile.logger)
library(nycflights13)
while (TRUE) {

  ## pick a ~car~flight
  flight <- flights[sample(1:nrow(flights), 1), ]

  ## prr <- .jnew('com.amazonaws.services.kinesis.model.PutRecordRequest')
  ## prr$setStreamName('test1')
  ## prr$setData(J('java.nio.ByteBuffer')$wrap(.jbyte(charToRaw(toJSON(car)))))
  ## prr$setPartitionKey(rownames(car))
  ## kc$putRecord(prr)

  res <- kinesis_put_record(stream = 'test-AWR', region = 'us-east-1',
                            data = toJSON(flight), partitionKey = flight$dest)
  flog.info(paste('Pushed a new flight to Kinesis:', res$sequenceNumber))

}
```

Write Data to the Stream from R



Minibuf-1

File Edit Options Buffers Tools Minibuf Help

```
library(rJava)
.jcall("java.util.Lang/System", "S", "setProperty", "aws.profile", "personal")
library(AWR.Kinesis); library(jsonlite); library(futile.logger); library(nycflights13)
while(TRUE) {
  flight <- flights[sample(1:nrow(flights), 1), ]
  res <- kinesis_put_record("test-AWR", region = "us-east-1", data = toJSON(flight),
                           partitionKey = flight$dest)
  flag.info(paste('Pushed a new flight to Kinesis:', res$sequenceNumber))
}
INFO [2017-03-06 21:47:16] Pushed a new flight to Kinesis: 49571082550613725758991014186133813750336290428328869938
INFO [2017-03-06 21:47:17] Pushed a new flight to Kinesis: 49571082550613725758991014186304346157376748486886405678
INFO [2017-03-06 21:47:17] Pushed a new flight to Kinesis: 4957108255061372575899101418661013052326445391882580018
INFO [2017-03-06 21:47:17] Pushed a new flight to Kinesis: 4957108255061372575899101418670049574481599664444407858
INFO [2017-03-06 21:47:18] Pushed a new flight to Kinesis: 49571082550613725758991014187125132922420286489026887730
INFO [2017-03-06 21:47:18] Pushed a new flight to Kinesis: 49571082550613725758991014187125132922420286489026887730
INFO [2017-03-06 21:47:19] Pushed a new flight to Kinesis: 49571082550613725758991014187125132922420286489026887730
INFO [2017-03-06 21:47:19] Pushed a new flight to Kinesis: 49571082550613725758991014187125132922420286489026887730
INFO [2017-03-06 21:47:19] Pushed a new flight to Kinesis: 49571082550613725758991014187125132922420286489026887730
INFO [2017-03-06 21:47:20] Pushed a new flight to Kinesis: 49571082550613725758991014188005348616637380517784895538
INFO [2017-03-06 21:47:20] Pushed a new flight to Kinesis: 49571082550613725758991014188005348616637380517784895538
INFO [2017-03-06 21:47:21] Pushed a new flight to Kinesis: 495710825506137257589910141880507262084201931439080294194
INFO [2017-03-06 21:47:21] Pushed a new flight to Kinesis: 49571082550613725758991014188400229242192365224945810
INFO [2017-03-06 21:47:21] Pushed a new flight to Kinesis: 49571082550613725758991014188400229242192365224945810
INFO [2017-03-06 21:47:22] Pushed a new flight to Kinesis: 49571082550613725758991014188400229242192365224945810
INFO [2017-03-06 21:47:22] Pushed a new flight to Kinesis: 49571082550613725758991014188400229242192365224945810
INFO [2017-03-06 21:47:23] Pushed a new flight to Kinesis: 4957108255061372575899101418979444132129388110904987454
INFO [2017-03-06 21:47:23] Pushed a new flight to Kinesis: 4957108255061372575899101418979444132129388110904987454
INFO [2017-03-06 21:47:23] Pushed a new flight to Kinesis: 49571082550613725758991014190015674357118885258062214962
INFO [2017-03-06 21:47:24] Pushed a new flight to Kinesis: 49571082550613725758991014190342048524814815135236882482
INFO [2017-03-06 21:47:24] Pushed a new flight to Kinesis: 49571082550613725758991014190545184066110072903306996786
INFO [2017-03-06 21:47:25] Pushed a new flight to Kinesis: 49571082550613725758991014190772462106627391577874
INFO [2017-03-06 21:47:25] Pushed a new flight to Kinesis: 495710825506137257589910141909658902513359818014746034
INFO [2017-03-06 21:47:25] Pushed a new flight to Kinesis: 495710825506137257589910141910883625734956644
INFO [2017-03-06 21:47:26] Pushed a new flight to Kinesis: 495710825506137257589910141910883625734956644
INFO [2017-03-06 21:47:26] Pushed a new flight to Kinesis: 495710825506137257589910141910883625734956644
INFO [2017-03-06 21:47:27] Pushed a new flight to Kinesis: 495710825506137257589910141910901316454989380989229962
INFO [2017-03-06 21:47:27] Pushed a new flight to Kinesis: 495710825506137257589910141910901316454989380989229962
INFO [2017-03-06 21:47:27] Pushed a new flight to Kinesis: 4957108255061372575899101419179651774223051686948700210
INFO [2017-03-06 21:47:28] Pushed a new flight to Kinesis: 49571082550613725758991014192373079905367392423621165106
INFO [2017-03-06 21:47:28] Pushed a new flight to Kinesis: 495710825506137257589910141925955222517648491767101490
INFO [2017-03-06 21:47:28] Pushed a new flight to Kinesis: 495710825506137257589910141925955222517648491767101490
INFO [2017-03-06 21:47:29] Pushed a new flight to Kinesis: 49571082550613725758991014192606209993679922
INFO [2017-03-06 21:47:29] Pushed a new flight to Kinesis: 49571082550613725758991014192606209993679922
INFO [2017-03-06 21:47:29] Pushed a new flight to Kinesis: 49571082550613725758991014192606209993679922
INFO [2017-03-06 21:47:30] Pushed a new flight to Kinesis: 4957108255061372575899101419320629795081871993731917106
INFO [2017-03-06 21:47:30] Pushed a new flight to Kinesis: 49571082550613725758991014193439653605558828482136114
INFO [2017-03-06 21:47:30] Pushed a new flight to Kinesis: 4957108255061372575899101419360611018833158206748536626
INFO [2017-03-06 21:47:30] Pushed a new flight to Kinesis: 4957108255061372575899101419360611018833158206748536626
INFO [2017-03-06 21:47:31] Pushed a new flight to Kinesis: 495710825506137257589910141940381538807481014923530546
INFO [2017-03-06 21:47:31] Pushed a new flight to Kinesis: 495710825506137257589910141940381538807481014923530546
INFO [2017-03-06 21:47:32] Pushed a new flight to Kinesis: 49571082550613725758991014194267446664703516545644173106
INFO [2017-03-06 21:47:32] Pushed a new flight to Kinesis: 49571082550613725758991014194267446664703516545644173106
INFO [2017-03-06 21:47:32] Pushed a new flight to Kinesis: 49571082550613725758991014194267446664703516545644173106
INFO [2017-03-06 21:47:33] Pushed a new flight to Kinesis: 4957108255061372575899101419516932321636030048312933874
INFO [2017-03-06 21:47:34] Pushed a new flight to Kinesis: 4957108255061372575899101419547881035957375117038714930
INFO [2017-03-06 21:47:34] Pushed a new flight to Kinesis: 4957108255061372575899101419547881035957375117038714930
INFO [2017-03-06 21:47:35] Pushed a new flight to Kinesis: 4957108255061372575899101419621625508592229898232895026
INFO [2017-03-06 21:47:35] Pushed a new flight to Kinesis: 4957108255061372575899101419621625508592229898232895026
INFO [2017-03-06 21:47:36] Pushed a new flight to Kinesis: 4957108255061372575899101419658613164183219466302644146
INFO [2017-03-06 21:47:36] Pushed a new flight to Kinesis: 4957108255061372575899101419658613164183219466302644146
INFO [2017-03-06 21:47:37] Pushed a new flight to Kinesis: 49571082550613725758991014197254722364971265580840517682
INFO [2017-03-06 21:47:37] Pushed a new flight to Kinesis: 49571082550613725758991014197254722364971265580840517682
INFO [2017-03-06 21:47:37] Pushed a new flight to Kinesis: 49571082550613725758991014197254722364971265580840517682
INFO [2017-03-06 21:47:38] Pushed a new flight to Kinesis: 49571082550613725758991014197254722364971265580840517682
INFO [2017-03-06 21:47:38] Pushed a new flight to Kinesis: 49571082550613725758991014197254722364971265580840517682
INFO [2017-03-06 21:47:38] Pushed a new flight to Kinesis: 49571082550613725758991014197254722364971265580840517682
INFO [2017-03-06 21:47:39] Pushed a new flight to Kinesis: 4957108255061372575899101419814287330016208053637464498
INFO [2017-03-06 21:47:39] Pushed a new flight to Kinesis: 4957108255061372575899101419814287330016208053637464498
INFO [2017-03-06 21:47:40] Pushed a new flight to Kinesis: 49571082550613725758991014198604015135424693088991702322
INFO [2017-03-06 21:47:40] Pushed a new flight to Kinesis: 49571082550613725758991014198604015135424693088991702322
INFO [2017-03-06 21:47:41] Pushed a new flight to Kinesis: 49571082550613725758991014199624216971459391068200632
INFO [2017-03-06 21:47:41] Pushed a new flight to Kinesis: 49571082550613725758991014199624216971459391068200632
* 12U: R-kinesis-putrecord iESS [R]: run
j-x butterfly
```

< 109137 : 0 99%

Reading Data from the Stream



CARD.COM

```
## get an iterator
sir <- .jnew('com.amazonaws.services.kinesis.model.GetShardIteratorRequest')
sir$setStreamName('test-AWR')
sir$setShardId(.jnew('java/lang/String', '0'))
sir$setShardIteratorType('TRIM_HORIZON')
kc <- .jnew('com.amazonaws.services.kinesis.AmazonKinesisClient')
kc$setEndpoint('kinesis.us-east-1.amazonaws.com')
iterator <- kc$getShardIterator(sir)$getShardIterator()

## get records
grr <- .jnew('com.amazonaws.services.kinesis.model.GetRecordsRequest')
grr$setShardIterator(iterator)
records <- kc$getRecords(grr)$getRecords()

## transform to string
json <- sapply(records, function(x)
  rawToChar(x$getData()$array()))

## decode JSON
json[1]
fromJSON(json[1])
rbindlist(lapply(json, fromJSON))
```

Running the MultiLangDaemon locally



CARD.COM

```
Terminix: Default
1 / 1 ▾ + ⌂ ↻
1: daroczig@gergely-CARD: ~/Projects/card-rocker/r-kinesis-example/files
~/Projects/card-rocker/r-kinesis-example/files | master ? export AWS_PROFILE=personal
~/Projects/card-rocker/r-kinesis-example/files | master ? /usr/bin/java -cp \
"/usr/local/lib/R/site-library/AWR/java":~/usr/local/lib/R/site-library/AWR.Kinesis/java/*:." \
com.amazonaws.services.kinesis.multilang.MultilangDaemon ./app.properties
Mar 05, 2017 5:32:33 PM com.amazonaws.services.kinesis.clientlibrary.config.KinesisClientLibConfigurator getConfiguration
INFO: Value of workerId is not provided in the properties. WorkerId is automatically assigned as:
Mar 05, 2017 5:32:34 PM com.amazonaws.services.kinesis.clientlibrary.config.KinesisClientLibConfigurator withProperty
INFO: Successfully set property regionName with value us-east-1
Mar 05, 2017 5:32:34 PM com.amazonaws.services.kinesis.multilang.MultilangDaemonConfig buildExecutorService
INFO: Using a cached thread pool.
Mar 05, 2017 5:32:34 PM com.amazonaws.services.kinesis.multilang.MultilangDaemonConfig <init>
INFO: Running AWR-demo-app to process stream test-AWR with executable /app/app.R
Mar 05, 2017 5:32:34 PM com.amazonaws.services.kinesis.multilang.MultilangDaemonConfig prepare
INFO: Using workerId:
Mar 05, 2017 5:32:34 PM com.amazonaws.services.kinesis.multilang.MultilangDaemonConfig prepare
INFO: Using credentials with access key id:
Mar 05, 2017 5:32:34 PM com.amazonaws.services.kinesis.multilang.MultilangDaemonConfig prepare
INFO: MultiLangDaemon is adding the following fields to the User Agent: amazon-kinesis-client-library-java-1.7.3 amazon-kinesis-multi-lang-dae
mon-1.0.1 R /app/app.R
Mar 05, 2017 5:32:34 PM com.amazonaws.services.kinesis.leases.impl.LeaseCoordinator <init>
INFO: With failover time 10000 ms and epsilon 25 ms, LeaseCoordinator will renew leases every 3308 ms, takeleases every 20050 ms, process maxi
mum of 2147483647 leases and steal 1 lease(s) at a time.
Mar 05, 2017 5:32:34 PM com.amazonaws.services.kinesis.clientlibrary.lib.worker.Worker initialize
INFO: Initialization attempt 1
Mar 05, 2017 5:32:34 PM com.amazonaws.services.kinesis.clientlibrary.lib.worker.Worker initialize
INFO: Initializing LeaseCoordinator
Mar 05, 2017 5:32:35 PM com.amazonaws.services.kinesis.clientlibrary.lib.worker.Worker initialize
INFO: Syncing Kinesis shard info
Mar 05, 2017 5:32:36 PM com.amazonaws.services.kinesis.clientlibrary.lib.worker.Worker initialize
INFO: Starting LeaseCoordinator
Mar 05, 2017 5:32:36 PM com.amazonaws.services.kinesis.leases.impl.LeaseTaker computeLeasesToTake
INFO: Worker [REDACTED] needed 2 leases but none were expired, so it will steal lease shardId-000000000002 from befs
4447-3adb-444f-8dc6-67504e5c86ef
Mar 05, 2017 5:32:36 PM com.amazonaws.services.kinesis.leases.impl.LeaseTaker computeLeasesToTake
INFO: Worker [REDACTED] saw 3 total leases, 0 available leases, 2 workers. Target is 2 leases, I have 0 leases, I wi
ll take 1 leases
Mar 05, 2017 5:32:36 PM com.amazonaws.services.kinesis.leases.impl.LeaseTaker takeLeases
INFO: Worker [REDACTED] successfully took 1 leases: shardId-000000000002
Mar 05, 2017 5:32:46 PM com.amazonaws.services.kinesis.clientlibrary.lib.worker.Worker run
INFO: Initialization complete. Starting worker loop.
Mar 05, 2017 5:32:46 PM com.amazonaws.services.kinesis.clientlibrary.lib.worker.Worker infoForce
INFO: Created new shardConsumer for : ShardInfo [shardId=shardId-000000000002, concurrencyToken=
shardIds=[shardId-000000000000], checkpoint={SequenceNumber: TRIM_HORIZON, SubsequenceNumber: 0}]
Mar 05, 2017 5:32:46 PM com.amazonaws.services.kinesis.clientlibrary.lib.worker.BlockOnParentShardTask call
INFO: No need to block on parents [shardId-000000000000] of shard shardId-000000000002
Mar 05, 2017 5:32:47 PM com.amazonaws.services.kinesis.clientlibrary.lib.worker.KinesisDataFetcher initialize
```

This Kinesis app is being run



CARD.COM

```
library(futile.logger)
library(AWR.Kinesis)

kinesis_consumer(
    initialize      = function()
        flog.info('Hello'),
    processRecords = function(records)
        flog.info(paste('Received', nrow(records), 'records from Kinesis')),
    shutdown       = function()
        flog.info('Bye'),
    updater        = list(
        list(1, function()
            flog.info('Updating some data every minute')),
        list(1/60*10, function()
            flog.info(paste(
                'This is a high frequency updater call',
                'running every 10 seconds')))),
    checkpointing = 1,
    logfile = '/logs/logger.log')
```

Running the MultiLangDaemon locally



CARD.COM

```
Terminix: Default
1: ec2-user@ip-10-10-10-10 logs]$ head -n 44 logger.log
INFO [2017-03-05 03:35:23] Starting R Kinesis Consumer application
INFO [2017-03-05 03:35:23 UTC] shardId-000000000000 Start of initialize
INFO [2017-03-05 03:35:23 UTC] shardId-000000000000 Hello
INFO [2017-03-05 03:35:23 UTC] shardId-000000000000 End of initialize
INFO [2017-03-05 03:35:23 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:24 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:25 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:26 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:27 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:28 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:29 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:30 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:31 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:32 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:33 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:33 UTC] shardId-000000000000 This is a high frequency updater call running every 10 seconds
INFO [2017-03-05 03:35:34 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:35 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:36 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:37 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:38 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:39 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:40 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:41 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:42 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:43 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:43 UTC] shardId-000000000000 This is a high frequency updater call running every 10 seconds
INFO [2017-03-05 03:35:44 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:45 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:46 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:47 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:48 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:49 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:50 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:51 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:52 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:53 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:54 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:54 UTC] shardId-000000000000 This is a high frequency updater call running every 10 seconds
INFO [2017-03-05 03:35:55 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:56 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:57 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:58 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:59 UTC] shardId-000000000000 Received 3 records from Kinesis
[ec2-user@ip-10-10-10-10 logs]$
```

① Dockerize your Kinesis Consumer:

- Java
- R
- AWR, AWR.Kinesis packages
- app.R
- app.properties
- startup command

② Put it on Docker Hub

③ Run as a EC2 Container Service Task:

- Create an ECS cluster
- Create ECS Task Role
- Create a Task definition
- Run it (as a service)

Dockerize your Kinesis Consumer



CARD.COM

card-rocker/Dockerfile

Github, Inc. [US] | https://github.com/cardcorp/card-rocker/blob/master/r-kinesis/Dockerfile

Features Explore Pricing This repository Search Sign in or Sign up

cardcorp / card-rocker

Code Issues 0 Pull requests 0 Projects 0 Pulse Graphs Watch 13 Star 6 Fork 2

Branch: master card-rocker / r-kinesis / Dockerfile

daroczig break long lines 3bf49ac 26 seconds ago

1 contributor

18 lines (14 sloc) | 606 Bytes

```
FROM cardcorp/r-aws-java-pandoc:latest
MAINTAINER Gergely Daroczi <gergely.daroczi@card.com>

## Install extra AWR packages
RUN install2_r --error \
    AWR.KMS \
    AWR.Kinesis \
    && rm -rf /tmp/downloaded_packages/ /tmp/*.rds

## Run MultiLangDaemon on /app
ENTRYPOINT ["/usr/bin/java", \
    "-cp", \
    "/usr/local/lib/R/site-library/AWR/java/*:/usr/local/lib/R/site-library/AWR.Kinesis/java/*", \
    "com.amazonaws.services.kinesis.multilang.MultiLangDaemon"]

## Override this if the consumer app is mounted elsewhere or the config file has a different name
CMD ["app/app.properties"]
```

Raw Blame History

© 2017 GitHub, Inc. Terms Privacy Security Status Help Contact GitHub API Training Shop Blog About

Dockerize your Kinesis Consumer

A screenshot of a GitHub repository page for "card-rocker/Dockerfile". The repository has 13 watchers, 6 stars, and 2 forks. The Dockerfile contains the following code:

```
FROM cardcorp/r-kinesis:latest
MAINTAINER Gergely Daroczi <gergely.daroczi@card.com>
## Add consumer
COPY files /app
```

Dockerize your Kinesis Consumer



The screenshot shows a GitHub repository page for `cardcorp/card-rocker`. The file `app.R` is displayed, which contains R code for a Kinesis consumer. The code uses the `futile.logger` and `AWR.Kinesis` packages to log messages and interact with Kinesis. It defines functions for initializing, processing records, shutting down, and updating data every minute.

```
Executable File | 24 lines (17 sloc) 619 Bytes
Raw Blame History

1 #!/usr/bin/r --vanilla
2 library(futile.logger)
3 library(AWR.Kinesis)
4
5 kinesis_consumer(
6
7   initialize      = function()
8     flog.info('Hello'),
9
10  processRecords = function(records)
11    flog.info(paste('Received', nrow(records), 'records from Kinesis')),
12
13  shutdown        = function()
14    flog.info('Bye'),
15
16  updater         = list(
17    list(1, function()
18      flog.info('Updating some data every minute')),
19    list(1/60*10, function()
20      flog.info('This is a high frequency updater call running every 10 seconds'))),
21
22  checkpointing = 1,
23  logfile       = '/logs/logger.log')
```

Dockerize your Kinesis Consumer

A screenshot of a GitHub repository page. The URL is https://github.com/cardcorp/card-rocker/blob/master/r-kinesis-example/files/app.properties. The repository name is cardcorp / card-rocker. The branch is master. The file contains the following code:

```
1 executableName = /app/app.R
2 streamName = test-AWR
3 applicationName = AWR-demo-app
4 AWS CredentialsProvider = DefaultAWSCredentialsProviderChain
5 processingLanguage = R
6 regionName = us-east-1
```

Put it on Docker Hub



A screenshot of a web browser displaying the Docker Hub interface for the repository `cardcorp/r-kinesis-example`. The page title is "PUBLIC | AUTOMATED BUILD". Below the title, it says "Last pushed: a day ago". There are tabs for "Repo Info", "Tags", "Dockerfile", and "Build Details", with "Build Details" being the active tab. A table lists nine build entries. The first entry is currently building ("Building"), while the others are marked as "Success". The table columns are "Status", "Actions", "Tag", "Created", and "Last Updated". To the right of the table, there is a sidebar titled "Source Repository" containing the URL `cardcorp/card-rocker`.

Status	Actions	Tag	Created	Last Updated
Building	<button>Cancel</button>	latest	3 minutes ago	a minute ago
✓ Success		latest	a day ago	a day ago
✓ Success		latest	a day ago	a day ago
✓ Success		latest	a day ago	a day ago
✓ Success		latest	a day ago	a day ago
✓ Success		latest	a day ago	a day ago
✓ Success		latest	a day ago	a day ago
✓ Success		latest	a day ago	a day ago

Create an ECS cluster

The screenshot shows the 'Create Cluster' wizard in the AWS Management Console. The left sidebar shows 'Amazon ECS' and 'Clusters' selected. The main form is titled 'Create Cluster' and contains the following fields:

- Cluster name***: AWR-test
- Create an empty cluster
- EC2 Instance type***: t2.medium
- Number of instances***: 1
- EC2 Ami Id***: amzn-ami-2016.09.f-amazon-ecs-optimized [ami-b2df2ca4]
- EBS storage (GiB)***: 22
- Key pair**: (dropdown menu)

A note below the key pair field states: "You will not be able to SSH into your EC2 instances without a key pair. You can create a new key pair in the EC2 console."

Networking

Configure the VPC for your container instances to use. A VPC is an isolated portion of the AWS cloud populated by AWS objects, such as Amazon EC2 instances. You can choose an existing VPC, or create a new one with this wizard.

VPC: Create a new vpc

CIDR Block: 10.0.0.0/16



Create ECS Task Role

IAM Management

https://console.aws.amazon.com/iam/home?region=us-east-1#roles

Gergely Daróczí Global Support

Create Role

Step 1 : Set Role Name

Step 2 : Select Role Type

Step 3 : Establish Trust

Step 4 : Attach Policy

Step 5 : Review

Set Role Name

Enter a role name. You cannot edit the role name after the role is created.

Role Name Maximum 64 characters. Use alphanumeric and '-' characters

Cancel Next Step

The screenshot shows the 'Create Role' wizard in the AWS IAM console. The user is on Step 1: Set Role Name. The role name 'AWR-ECS' has been entered into the input field. The 'Next Step' button is visible at the bottom right. The sidebar on the left lists the remaining steps: Step 2: Select Role Type, Step 3: Establish Trust, Step 4: Attach Policy, and Step 5: Review.

Create ECS Task Role

The screenshot shows the AWS IAM Management console with the URL <https://console.aws.amazon.com/iam/home?region=us-east-1#roles>. The left sidebar shows the navigation path: IAM Management > Services > Resource Groups. The main area is titled "Create Role" and "Step 2 : Select Role Type". A sub-section titled "AWS Service Roles" is expanded, showing the following options:

- Amazon EC2 Role for EC2 Container Service**: Role to allow EC2 instances in an Amazon ECS cluster to access Amazon ECS.
- Amazon EC2 Container Service Role**: Allows ECS to create and manage AWS resources on your behalf.
- Amazon EC2 Container Service Task Role**: Allows ECS tasks to call AWS services on your behalf.
- Amazon EC2 Spot Fleet Role**: Role to Allow EC2 Spot Fleet to request and terminate Spot Instances on your behalf.
- Amazon Elastic MapReduce**: Role to allow EMR to access other AWS services such as EC2 on your behalf.
- Role for Cross-Account Access** (radio button selected):
- Role for Identity Provider Access** (radio button unselected):

At the bottom right of the main area are "Cancel", "Previous", and "Next Step" buttons.

Create ECS Task Role



IAM Management > Services > Resource Groups

https://console.aws.amazon.com/iam/home?region=us-east-1#roles

Create Role

Step 1 : Set Role Name

Step 2 : Select Role Type

Step 3 : Establish Trust

Step 4 : Attach Policy

Step 5 : Review

Attach Policy

Select one or more policies to attach. Each role can have up to 10 policies attached.

Filter: Policy Type ▾ kinesis Showing 7 results

	Policy Name	Attached Entities	Creation Time	Edited Time
<input type="checkbox"/>	AmazonKinesisFullAccess	1	2015-02-06 10:40 PST	2015-02-06 10:40 PST
<input checked="" type="checkbox"/>	AmazonKinesisReadOnlyA...	1	2015-02-06 10:40 PST	2015-02-06 10:40 PST
<input type="checkbox"/>	AmazonKinesisAnalyticsF...	0	2016-09-21 12:01 PST	2016-09-21 12:01 PST
<input type="checkbox"/>	AmazonKinesisAnalyticsR...	0	2016-09-21 11:16 PST	2016-09-21 11:16 PST
<input type="checkbox"/>	AmazonKinesisFirehoseFu...	0	2015-10-07 11:45 PST	2015-10-07 11:45 PST
<input type="checkbox"/>	AmazonKinesisFirehoseR...	0	2015-10-07 11:43 PST	2015-10-07 11:43 PST
<input type="checkbox"/>	AWSLambdaKinesisExecu...	0	2015-04-09 08:14 PST	2015-04-09 08:14 PST

Cancel Previous Next Step

Create a Task definition



Amazon EC2 Cont... https://console.aws.amazon.com/ecs/home?region=us-east-1#/taskDefinitions/create Gergely Darócz N. Virginia Support

Create a Task Definition

A task definition specifies which containers are included in your task and how they interact with each other. You can also specify data volumes for your containers to use. [Learn more](#)

Task Definition Name* AWR-logger

Task Role AWR-ECS

Optional IAM role that tasks can use to make API requests to authorized AWS services. Create an Amazon EC2 Container Service Task Role in the [IAM Console](#).

Network Mode Bridge

Constraint

Constraints allow you to filter the instances used for your placement strategies using built-in or custom attributes. The scheduler first filters the instances that match the constraints and then applies the placement strategy to place the task.

Type	Expression
Add constraint	

Container Definitions

Add container

Container Name	Image	Hard/Soft memory limits (MB)	Essential
No results			

Volumes

Name	Source Path
No results	

Add volume

Configure via JSON

Create a Task definition



Amazon EC2 Cont... https://console.aws.amazon.com/ecs/home?region=us-east-1#/taskDefinitions/create Gergely Daroczi N. Virginia Support

Services Resource Groups

Amazon ECS Clusters Task Definitions Repositories

Create a Task Definition

A task definition is a template for your tasks. You can reuse it for other tasks or services. You can also specify data volumes.

Add volume

Name* logs

Source path /logs

*Required Cancel Add

Constraint

Constraints allow you to filter the instances used for your placement strategies using built-in or custom attributes. The scheduler first filters the instances that match the constraints and then applies the placement strategy to place the task.

Type	Expression
<input checked="" type="radio"/> Add constraint	

Container Definitions

Add container

Container Name	Image	Hard/Soft memory limits (MB)	Essential
No results			

Volumes

Name	Source Path
No results	

Add volume

Configure via JSON

A screenshot of the AWS CloudWatch Metrics console. The main navigation bar shows 'Amazon EC2 Cont...' and the URL 'https://console.aws.amazon.com/ecs/home?region=us-east-1#/taskDefinitions/create'. On the left, there's a sidebar with 'Amazon ECS' selected, followed by 'Clusters', 'Task Definitions' (which is highlighted in orange), and 'Repositories'. The main content area has a title 'Create a Task Definition' and a sub-section 'Add volume'. Inside this sub-section, there are two input fields: 'Name*' containing 'logs' and 'Source path' containing '/logs'. Below these fields is a note '*Required' and a row with 'Cancel' and 'Add' buttons. Underneath this, there's a section titled 'Constraint' with a note about filtering instances. Below that is a table header for 'Container Definitions' with a 'Type' column and an 'Add constraint' button. The table body shows a single row with columns for 'Container Name', 'Image', 'Hard/Soft memory limits (MB)', and 'Essential', all of which are currently empty. A note 'No results' is displayed below the table. Below the container definitions is another section titled 'Volumes' with a similar table structure, also showing 'No results'. At the bottom of the volume section are 'Add volume' and 'Configure via JSON' buttons.



Create a Task definition

The screenshot shows the 'Create a Task Definition' wizard in the AWS ECS console. The left sidebar lists 'Amazon ECS', 'Clusters', 'Task Definitions' (which is selected), and 'Repositories'. The main area is titled 'Add container' under 'Standard' type.

- Container name***: logger
- Image***: cardcorp/r-kinesis-example:latest
- Memory Limits (MB)***: Hard limit 512
- Port mappings**: Host port, Container port, Protocol (tcp)
- Advanced container configuration**:
 - ENVIRONMENT**: CPU units (input field), Essential (checkbox checked). A tooltip indicates: "The number of cpu units to reserve for the container. A container instance has 1.024 cpu units for every CPU core."

* Required

Cancel Add

Create a Task definition



CARD.COM

The screenshot shows the 'Create a Task Definition' wizard in the AWS ECS console. The left sidebar lists 'Amazon ECS', 'Clusters', 'Task Definitions' (which is selected), and 'Repositories'. The main area is titled 'Create a Task Definition' with the sub-instruction: 'A task definition specifies the resources required for your containers to run.' Below this, there's a 'Constraint' section and a 'Type' section with an 'Add constraint' button.

The central part of the screen is the 'Add container' configuration. It includes fields for 'IP address' and 'Add extra host'. The 'STORAGE AND LOGGING' section contains:

- 'Read only root file system': An unchecked checkbox.
- 'Mount points': A table with one entry: 'Source volume' (dropdown set to 'logs'), 'Container path' (input field containing '/logs'), and 'Read only' (checkbox checked).
- 'Add mount point': A blue 'Add mount point' button.
- 'Volumes from': A table with one entry: 'Source container' (dropdown set to '<none>') and 'Read only' (checkbox checked).
- 'Add volumes': A blue 'Add volumes' button.
- 'Log configuration': Fields for 'Log driver' (dropdown set to '<none>'), 'Log options', 'Key' (button 'Add key'), and 'Value' (button 'Add value').

The 'SECURITY' section is partially visible at the bottom. At the very bottom, there's a note '* Required' and a row with 'Cancel' and 'Add' buttons.

Run the ECS Task

The screenshot shows the AWS ECS Task Definitions console. The URL is <https://console.aws.amazon.com/ecs/home?region=us-east-1#/taskDefinitions/AWR-logger/status/ACTIVE>. The user is Gergely Daróczki from the N. Virginia region.

The left sidebar shows the navigation menu: Services, Resource Groups, Amazon ECS, Clusters, Task Definitions (which is selected and highlighted in orange), and Repositories.

The main content area displays the "Task Definition Name : AWR-logger". Below it, a message says "Select a revision for more details".

A dropdown menu is open under the "Actions" button, showing the following options: Create new revision, Run Task (which is highlighted with a black background), Create Service, Update Service, Deregister, and Dereject.

Below the dropdown, there is a table with one row:

Task Definition Name	Status
AWR-logger:1	Active

At the bottom of the page, there are links for Feedback, English, and a footer with copyright information: © 2008-2017, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use.

Run the ECS Task

The screenshot shows the AWS CloudWatch Metrics (ECS) interface for a cluster named 'AWR'. The left sidebar has 'Clusters' selected. The main area displays the following information:

- Status: ACTIVE
- Registered container instances: 1
- Pending tasks count: 0
- Running tasks count: 1

Below this, there are tabs for Services, Tasks, ECS Instances, and Metrics, with 'Tasks' selected. A 'Run new Task' button is visible. The 'Last updated on March 5, 2017 6:05:01 PM (0m ago)' message is displayed. The 'Desired task status' is set to 'Running'. A table lists one task:

Task	Task Definition	Group	Container Inst...	Last status	Desired status	Started By
0c9f224a-7808...	AWR-logger:1	family:AWR-log...	27b4935-70e2...	RUNNING	RUNNING	

Run the ECS Task

Terminix: Default

```

1: ec2-user@[REDACTED] ~
/ ssh -i ~/.ssh/[REDACTED] ec2-user@[REDACTED]
Last login: Mon Mar  6 02:05:29 2017 from [REDACTED]

[REDACTED]
|_ \_|_/_| Amazon ECS-Optimized Amazon Linux AMI 2016.09.f
 \_\_|_/_/ [REDACTED]

For documentation visit, http://aws.amazon.com/documentation/ecs
4 package(s) needed for security, out of 6 available
Run "sudo yum update" to apply all updates.
[ec2-user@[REDACTED] ~]$ docker ps
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS              PORTS          NAMES
AMES                cardcorp/r-kinesis-example:latest   "/usr/bin/java -cp /u"   22 hours ago      Up 22 hours         e
cs-AWR-logger-1-logger-92ca888bd5d681e59d01   amazon/amazon-ecs-agent:latest    "/agent"           23 hours ago      Up 23 hours         e
Zed65c5a3752        cs-agent
[ec2-user@[REDACTED] ~]$ pgrep app.R
29435
[ec2-user@[REDACTED] ~]$ head -n 10 /logs/logger.log
INFO [2017-03-05 03:35:23] Starting R Kinesis Consumer application
INFO [2017-03-05 03:35:23 UTC] shardId-000000000000 Start of initialize
INFO [2017-03-05 03:35:23 UTC] shardId-000000000000 Hello
INFO [2017-03-05 03:35:23 UTC] shardId-000000000000 End of initialize
INFO [2017-03-05 03:35:23 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:24 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:25 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:26 UTC] shardId-000000000000 Received 2 records from Kinesis
INFO [2017-03-05 03:35:27 UTC] shardId-000000000000 Received 3 records from Kinesis
INFO [2017-03-05 03:35:28 UTC] shardId-000000000000 Received 2 records from Kinesis
[ec2-user@[REDACTED] ~]$
```

Scaling the Kinesis Consumer up

```
Terminix: Default
1: ec2-user@ip-172-31-10-15: ~ %
INFO [2017-03-05 03:43:01 UTC] shardId-000000000000 Received 1 records from Kinesis
INFO [2017-03-05 03:43:01 UTC] shardId-000000000000 This is a high frequency updater call running every 10 seconds
INFO [2017-03-05 03:43:09 UTC] shardId-000000000000 Shutting down
INFO [2017-03-05 03:43:09 UTC] shardId-000000000000 Bye
INFO [2017-03-05 03:43:15] Starting R Kinesis Consumer application
INFO [2017-03-05 03:43:15 UTC] shardId-000000000002 Start of initialize
INFO [2017-03-05 03:43:15 UTC] shardId-000000000002 Hello
INFO [2017-03-05 03:43:15 UTC] shardId-000000000002 End of initialize
INFO [2017-03-05 03:43:15] Starting R Kinesis Consumer application
INFO [2017-03-05 03:43:16 UTC] shardId-000000000001 Start of initialize
INFO [2017-03-05 03:43:16 UTC] shardId-000000000001 Hello
INFO [2017-03-05 03:43:16 UTC] shardId-000000000001 End of initialize
INFO [2017-03-05 03:44:32 UTC] shardId-000000000002 Received 1 records from Kinesis
INFO [2017-03-05 03:44:32 UTC] shardId-000000000002 Updating some data every minute
INFO [2017-03-05 03:44:32 UTC] shardId-000000000002 This is a high frequency updater call running every 10 seconds
INFO [2017-03-05 03:44:33 UTC] shardId-000000000002 Received 3 records from Kinesis
INFO [2017-03-05 03:44:34 UTC] shardId-000000000002 Received 1 records from Kinesis
INFO [2017-03-05 03:44:35 UTC] shardId-000000000002 Received 1 records from Kinesis
INFO [2017-03-05 03:44:36 UTC] shardId-000000000001 Received 2 records from Kinesis
INFO [2017-03-05 03:44:36 UTC] shardId-000000000002 Received 1 records from Kinesis
INFO [2017-03-05 03:44:36 UTC] shardId-000000000001 Updating some data every minute
INFO [2017-03-05 03:44:36 UTC] shardId-000000000001 This is a high frequency updater call running every 10 seconds
INFO [2017-03-05 03:44:37 UTC] shardId-000000000001 Received 1 records from Kinesis
INFO [2017-03-05 03:44:38 UTC] shardId-000000000001 Received 1 records from Kinesis
INFO [2017-03-05 03:44:39 UTC] shardId-000000000002 Received 1 records from Kinesis
INFO [2017-03-05 03:44:39 UTC] shardId-000000000001 Received 2 records from Kinesis
INFO [2017-03-05 03:44:40 UTC] shardId-000000000001 Received 1 records from Kinesis
INFO [2017-03-05 03:44:40 UTC] shardId-000000000002 Received 1 records from Kinesis
INFO [2017-03-05 03:44:41 UTC] shardId-000000000001 Received 2 records from Kinesis
INFO [2017-03-05 03:44:42 UTC] shardId-000000000001 Received 2 records from Kinesis
INFO [2017-03-05 03:44:43 UTC] shardId-000000000002 Received 2 records from Kinesis
INFO [2017-03-05 03:44:43 UTC] shardId-000000000002 This is a high frequency updater call running every 10 seconds
INFO [2017-03-05 03:44:44 UTC] shardId-000000000002 Received 1 Records from Kinesis
INFO [2017-03-05 03:44:45 UTC] shardId-000000000002 Received 2 records from Kinesis
INFO [2017-03-05 03:44:45 UTC] shardId-000000000001 Received 1 records from Kinesis
INFO [2017-03-05 03:44:46 UTC] shardId-000000000002 Received 1 records from Kinesis
INFO [2017-03-05 03:44:46 UTC] shardId-000000000001 Received 1 records from Kinesis
INFO [2017-03-05 03:44:47 UTC] shardId-000000000001 Received 1 records from Kinesis
INFO [2017-03-05 03:44:47 UTC] shardId-000000000001 This is a high frequency updater call running every 10 seconds
INFO [2017-03-05 03:44:47 UTC] shardId-000000000002 Received 1 records from Kinesis
INFO [2017-03-05 03:44:48 UTC] shardId-000000000001 Received 1 records from Kinesis
INFO [2017-03-05 03:44:48 UTC] shardId-000000000002 Received 1 records from Kinesis
INFO [2017-03-05 03:44:49 UTC] shardId-000000000001 Received 1 records from Kinesis
INFO [2017-03-05 03:44:49 UTC] shardId-000000000002 Received 1 records from Kinesis
INFO [2017-03-05 03:44:50 UTC] shardId-000000000001 Received 1 records from Kinesis
```

Nice example project, but ...

- I might want to avoid publishing my Consumer on Docker Hub
- I might want to avoid publishing my code on GitHub
- I might want to avoid committing credentials etc to the repo

Problems:

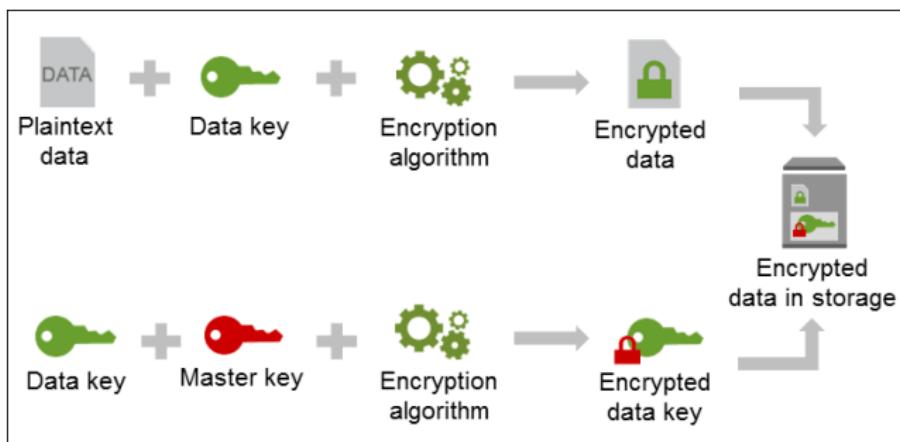
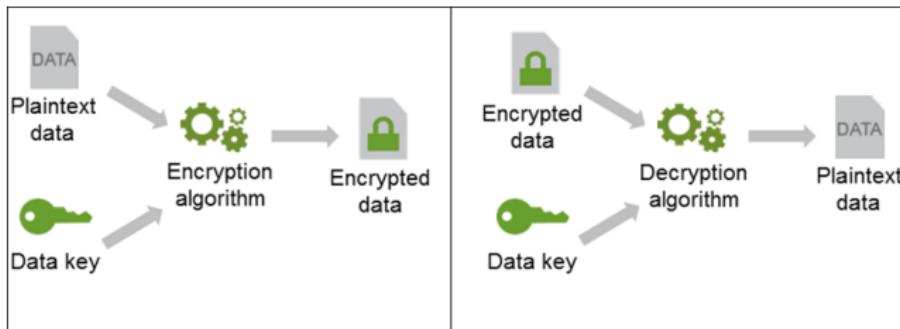
- How to store credentials in the Docker images?
- Where to store the Docker images?

Nice example project, but ...

- I might want to avoid publishing my Consumer on Docker Hub
- I might want to avoid publishing my code on GitHub
- I might want to avoid committing credentials etc to the repo

Problems:

- How to store credentials in the Docker images? **KMS**
- Where to store the Docker images? **ECR**



Source: AWS Encryption SDK

- encrypt up to 4 KB of arbitrary data:

```
> library(AWR.KMS)
> kms_encrypt('alias/mykey', 'foobar')
[1] "Base-64 encoded ciphertext"
```

- decrypt such Base-64 encoded ciphertext back to plaintext:

```
> kms_decrypt('Base-64 encoded ciphertext')
[1] "foobar"
```

- generate a data encryption key:

```
> kms_generate_data_key('alias/mykey')
$cipher
[1] "Base-64 encoded, encrypted data encryption key"
$key
[1] "alias/mykey"
$text
[1] 00 01 10 11 00 01 10 11 ...
```

Encrypting Data Larger Than 4 KB?



CARD.COM

```
## let's say we want to encrypt the mtcars dataset stored in JSON
library(jsonlite)
data <- toJSON(mtcars)

## generate a 256-bit data encryption key (that's supported by digest::AES)
library(AWR.KMS)
key <- kms_generate_data_key('alias/mykey', byte = 32L)

## convert the JSON to raw so that we can use that with digest::AES
raw <- charToRaw(data)
## the text length must be a multiple of 16 bytes
## https://github.com/sdoyen/r_password_crypt/blob/master/crypt.R
raw <- c(raw, as.raw(rep(0, 16 - length(raw) %% 16)))

## encrypt the raw object with the new key + digest::AES
## the resulting text and the encrypted key can be stored on disk
library(digest)
aes <- AES(key$text)
base64_enc(aes$encrypt(raw))

## decrypt the above returned ciphertext using the decrypted key
rawToChar(aes$decrypt(base64_dec(...), raw = TRUE))
```

Example “Production” Consumer App



CARD.COM

```
library(AWR.Kinesis); library(jsonlite); library(AWR.KMS); library(futile.logger); flog.threshold(DEBUG)

kinesis_consumer(
  initialize      = function() {
    flog.info('Decrypting Redis hostname via KMS')
    host <- kms_decrypt('AQECAHiiz4GEPFQLL9AAON5TY/lDR5euQQScpXQU9iYTn+u... ')
    flog.info('Connecting to Redis')
    library(rredis); redisConnect(host = host)
    flog.info('Connected to Redis')
  },
  processRecords = function(records) {
    flog.info(paste('Received', nrow(records), 'records from Kinesis'))
    for (record in records$data) {
      flight <- fromJSON(record)$dest
      if (!is.null(flight)) {
        flog.debug(paste('Adding +1 to', flight))
        redisIncr(sprintf('flight:%s', flight))
      } else {
        flog.error('Flight destination not found')
      }
    }
  },
  updaters = list(
    list(1/6, function() {
      flog.info('Checking overall counters')
      flights <- redisKeys('flight:*')
      for (flight in flights) {
        flog.debug(paste('Found', redisGet(flight), sub('^flight:', '', flight)))
      }
    })),
  logfile = '/logs/redis.log')
```

Dockerfile:

```
FROM cardcorp/r-kinesis:latest
MAINTAINER Gergely Daroczi <gergely.daroczi@card.com>

## Install R package to interact with Redis
RUN install2.r --error rredis && rm -rf /tmp/downloaded_packages/ /tmp/*.rds

## Add consumer
COPY files /app
```

Build and push to ECR:

```
docker build -t cardcorp/r-kinesis-secret .
`aws ecr get-login --region us-east-1`
docker tag -f cardcorp/r-kinesis-secret:latest \
    ***.dkr.ecr.us-east-1.amazonaws.com/cardcorp/r-kinesis-secret:latest
docker push ***.dkr.ecr.us-east-1.amazonaws.com/cardcorp/r-kinesis-secret:latest
```

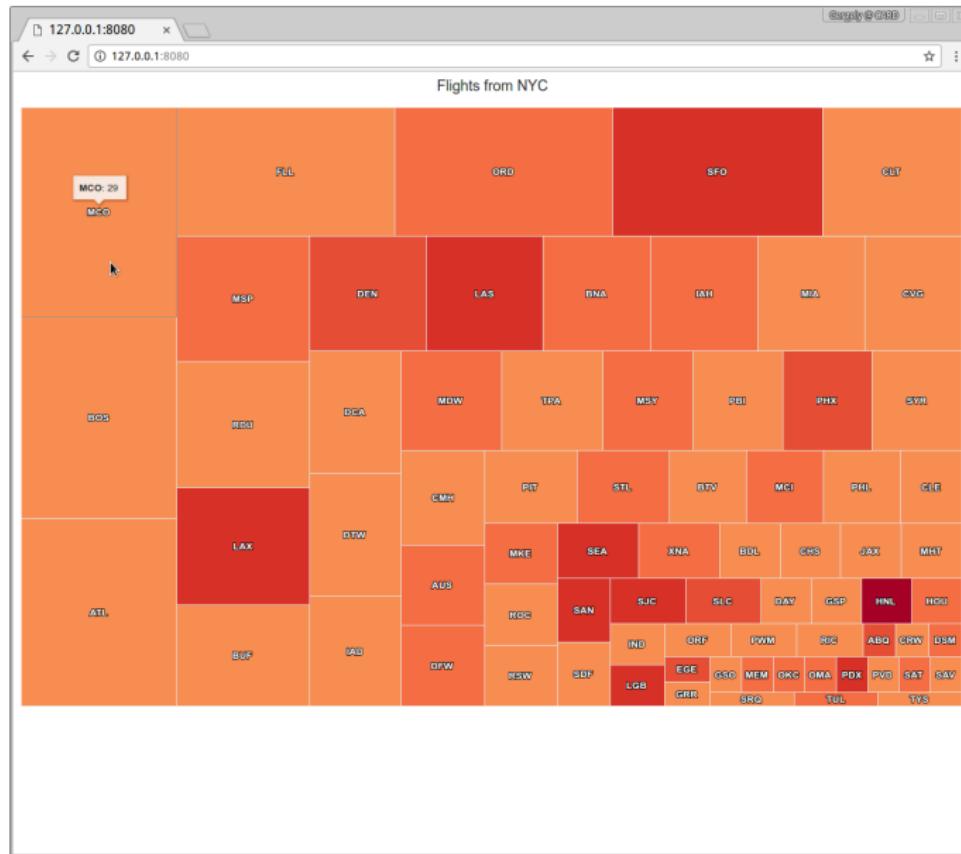
```
library(treemap);library(highcharter);library(nycflights13)
library(rredis);redisConnect(host = '***', port = '***')

ui      <- shinyUI(highchartOutput('treemap', height = '800px'))
server <- shinyServer(function(input, output, session) {

  destinations <- reactive({
    reactiveTimer(2000)()
    flights <- redisMGet(redisKeys('flight:*'))
    flights <- data.frame(faa = sub('^flight:', '', names(flights)),
                           N   = as.numeric(flights))
    merge(flights, airports, by = 'faa')
  })

  output$treemap <- renderHighchart({
    tm <- treemap(destinations(), index = c('faa'),
                  vSize = 'N', vColor = 'tz',
                  type = 'value', draw = FALSE)
    hc_title(hctreemap(tm, animation = FALSE), text = 'Flights from NYC')
  })
}

shinyApp(ui = ui, server = server)
```



- AWR repo:
 - 6.3 GB
 - 273 tags/versions
 - GitLab + CI + drat

```
install.packages('AWR', repos = 'https://cardcorp.gitlab.io/AWR')
```

- AWR repo:
 - 6.3 GB
 - 273 tags/versions
 - GitLab + CI + drat

```
install.packages('AWR', repos = 'https://cardcorp.gitlab.io/AWR')
```

- Submitted to CRAN on
 - 2016-12-05

- AWR repo:
 - 6.3 GB
 - 273 tags/versions
 - GitLab + CI + drat

```
install.packages('AWR', repos = 'https://cardcorp.gitlab.io/AWR')
```

- Submitted to CRAN on
 - 2016-12-05
 - 2017-01-09
 - 2017-01-10
 - 2017-01-11
 - 2017-01-11
 - 2017-01-13

- AWR repo:
 - 6.3 GB
 - 273 tags/versions
 - GitLab + CI + drat

```
install.packages('AWR', repos = 'https://cardcorp.gitlab.io/AWR')
```

- Submitted to CRAN on
 - 2016-12-05
 - 2017-01-09
 - 2017-01-10
 - 2017-01-11
 - 2017-01-11
 - 2017-01-13
- Release cycle: 2 minor, ~125 patch versions in the past 12 months
- CI



```
> library(rJava)
> kc <- .jnew('com.amazonaws.services.s3.AmazonS3Client')
> kc$getS3AccountOwner()$getDisplayName()
[1] "foobar"
```



Because "S"
is so 1992.